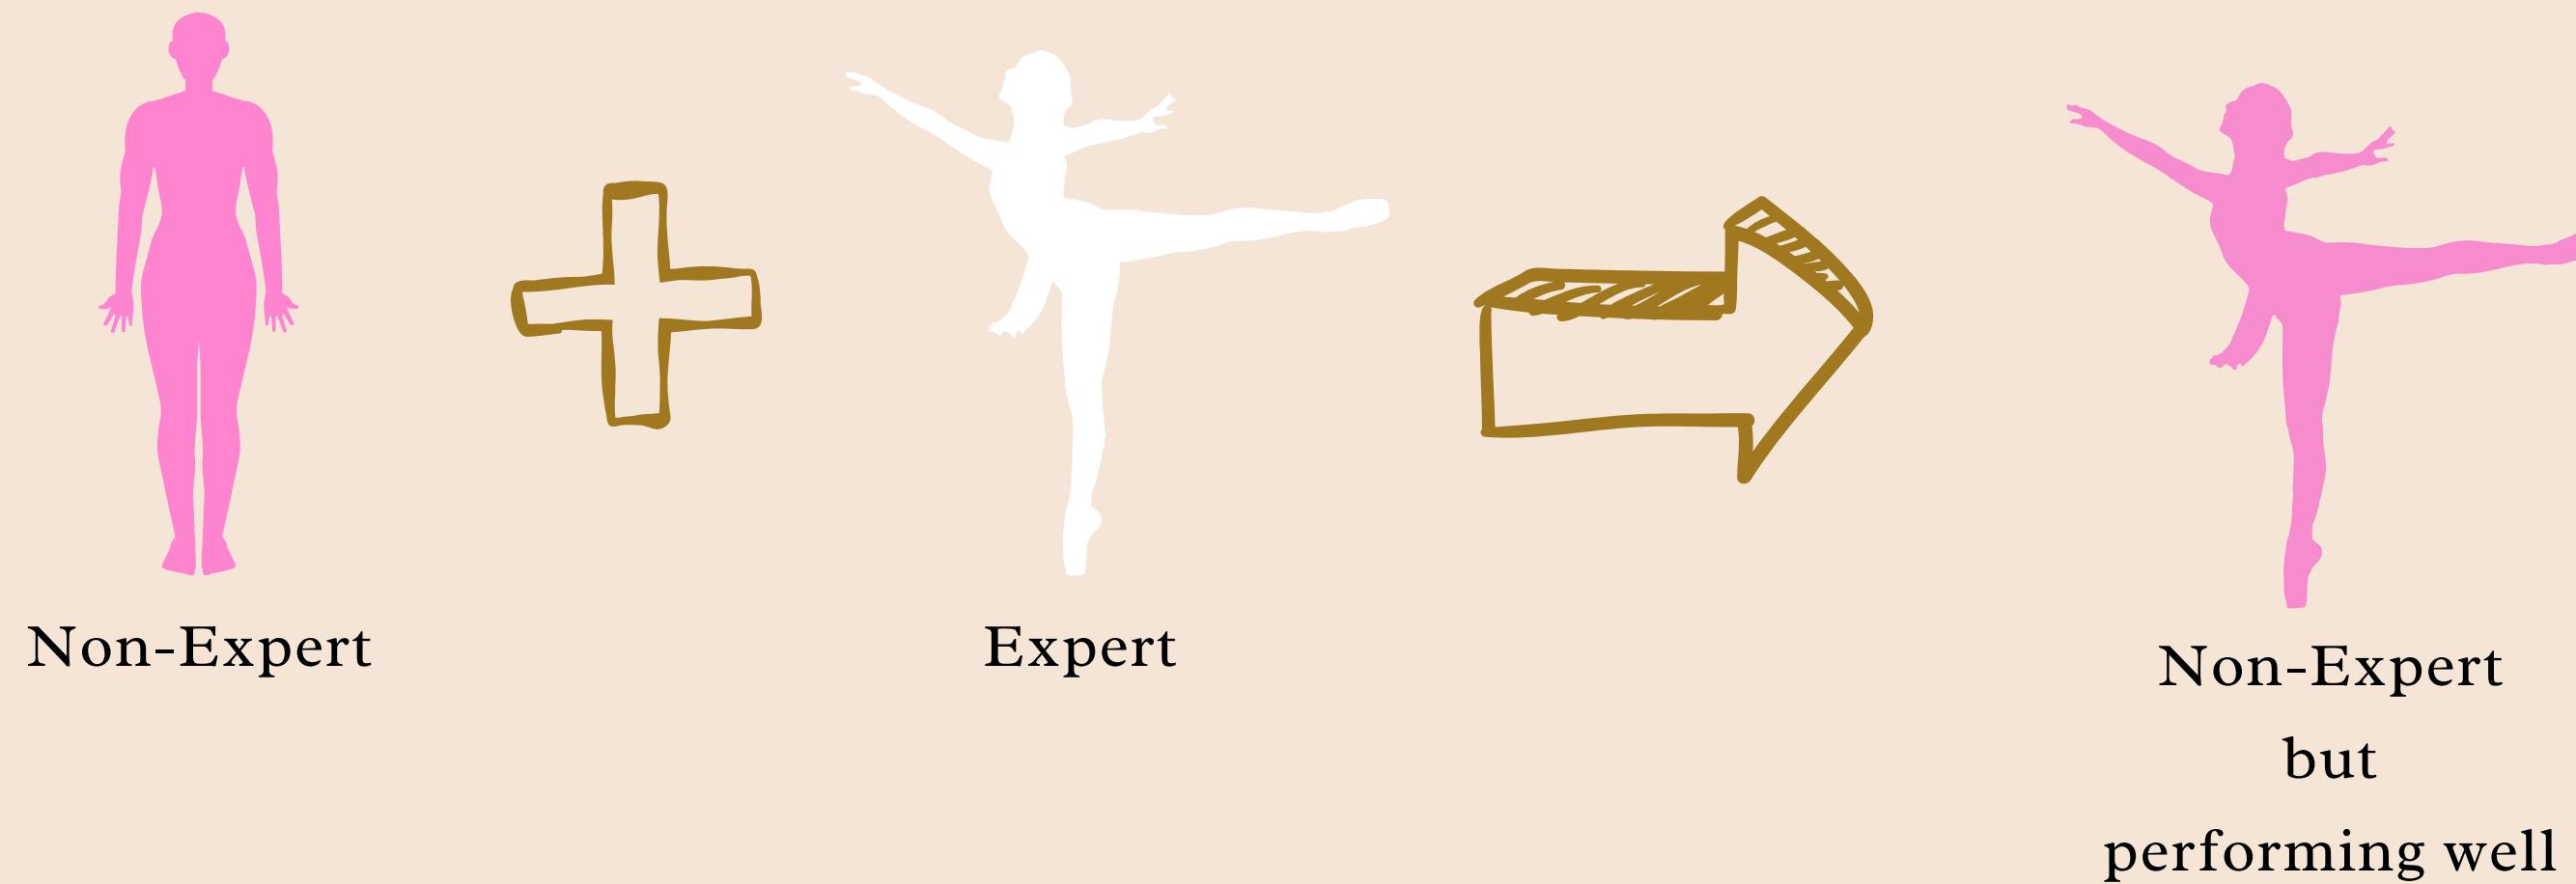


# Creating Illusions: Synthesizing Dynamic Movements for Non-Experts

AI 2023 Final Project  
Group 28 : 110704031 賴怡 110263029 楊芊華 109612007 吳宜靜

# Introduction

Developing an AI model to convert human motion to skeleton and generate an image of another human performing the same motion.



# Related Work

-Image-to-Image Translation with  
Conditional Adversarial Networks

Phillip Isola, Jun-Yan Zhu, Tinghui Zhou, Alexei A. Efros

CVPR, 2017

-PatchedGAN

Introduced by Isola et al.

-Keypoint R-CNN

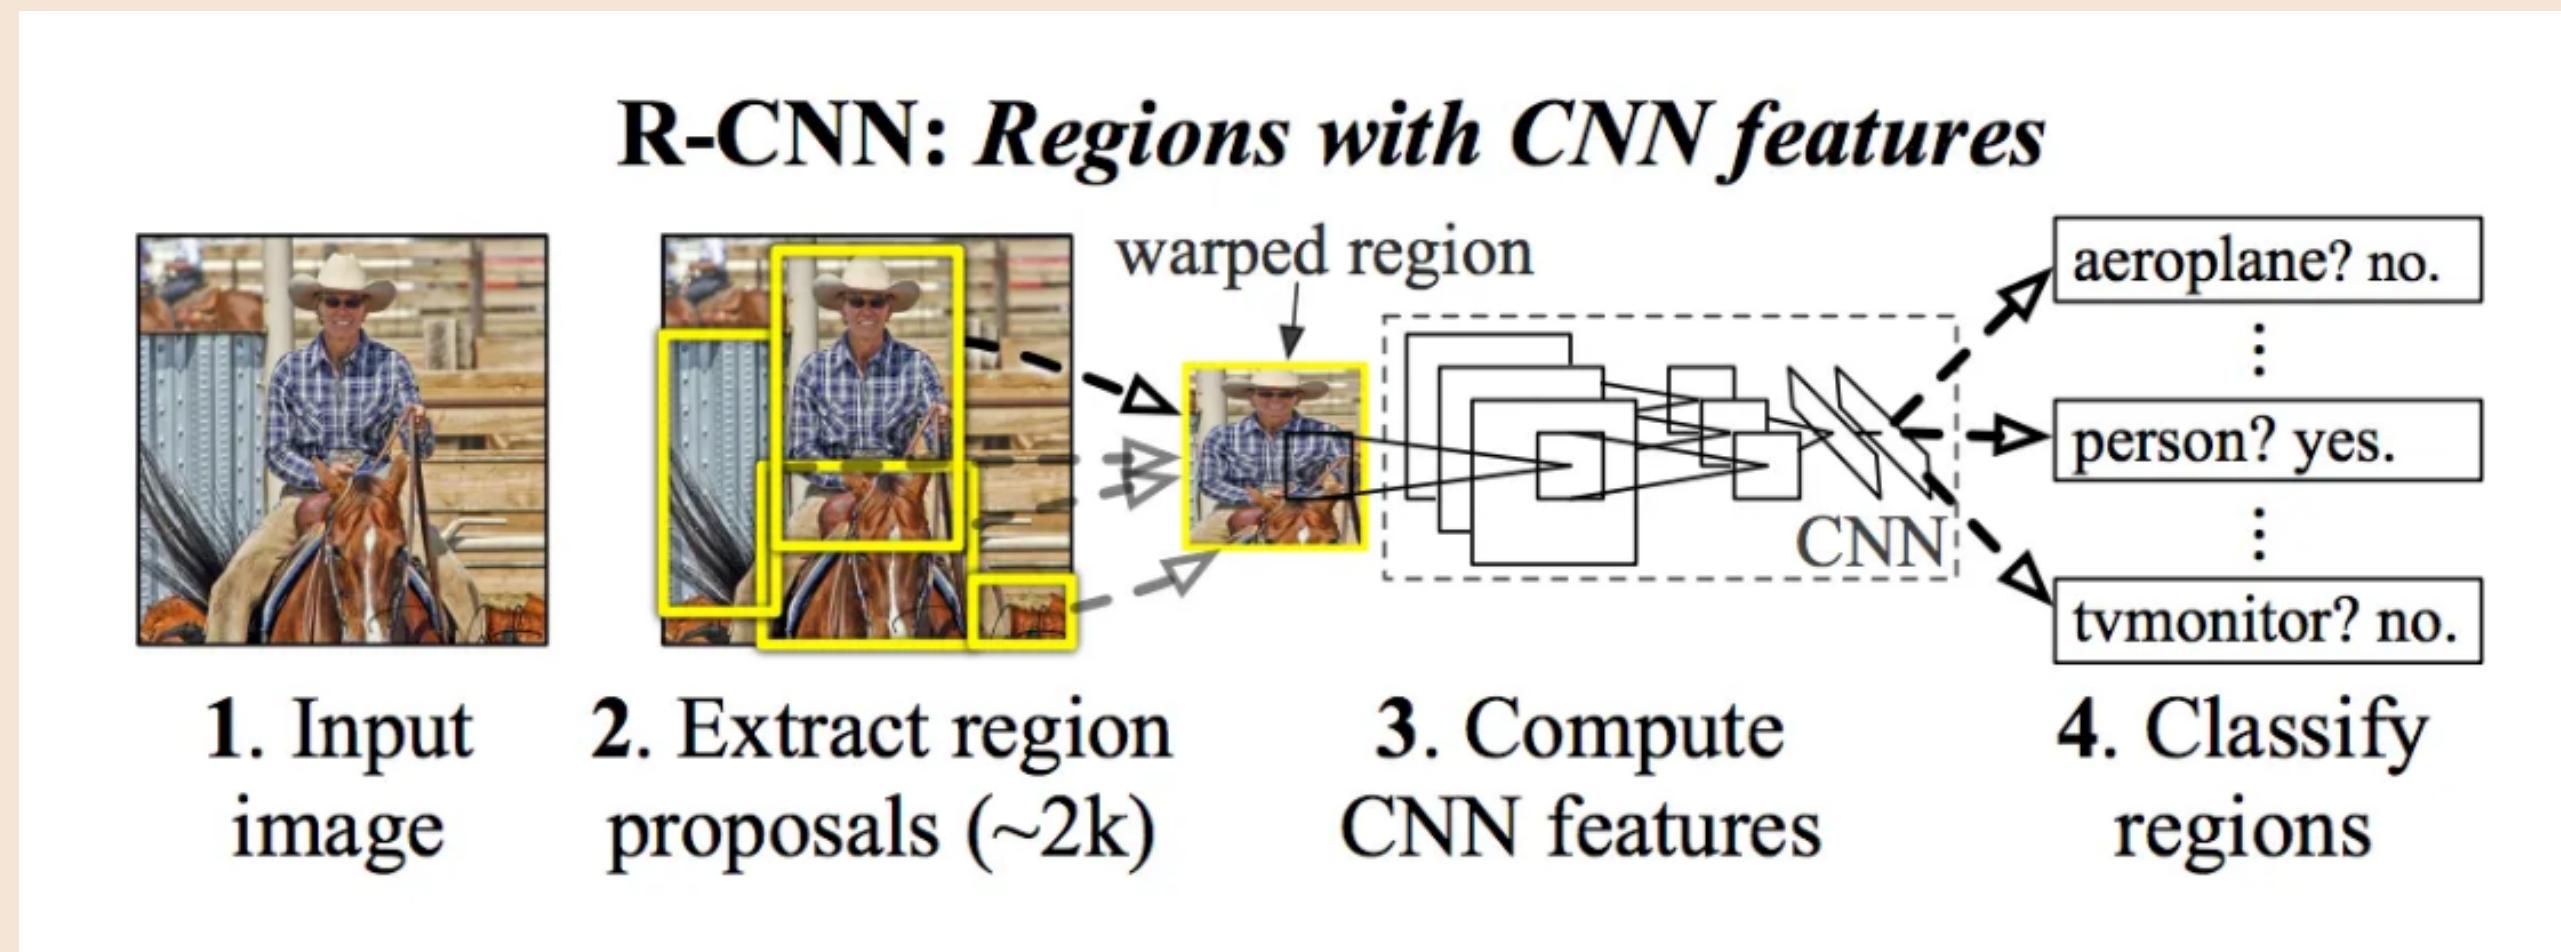
Kaiming He Georgia Gkioxari Piotr Dollar Ross Girshick  
Facebook AI Research (FAIR)

# Dataset

- Keypoint Detection - COCO
  - pros:
    - wide range of coverage
    - high-quality annotations
  - dataset selection:
    - get 100,000+ keypoint data
    - remove data with less than 13 keypoints and randomly select 10,000 samples
- Input Video - dancing videos on youtube
  - video segmentation: extract 2,500+ images

# Baseline - image to skeleton

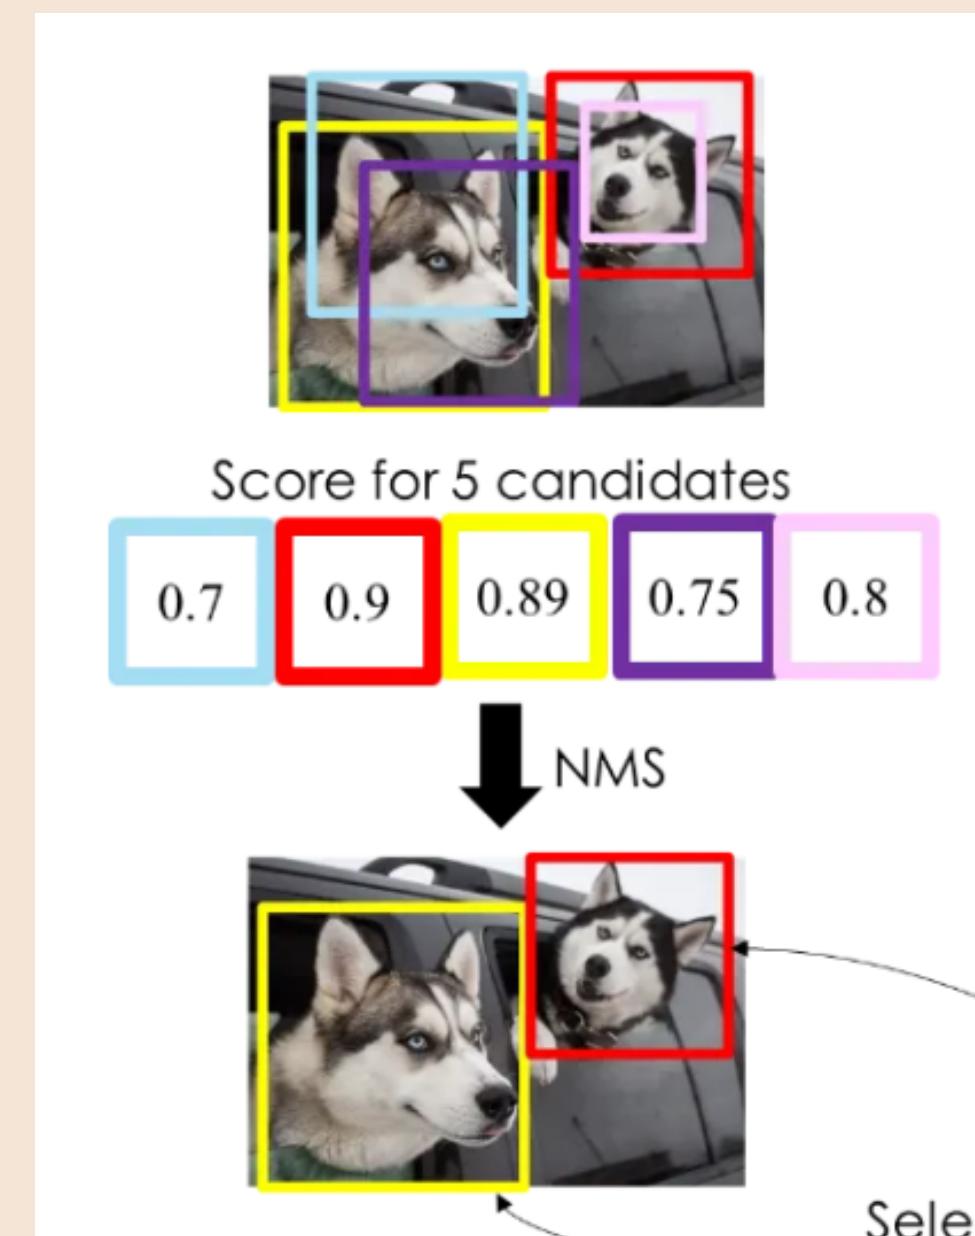
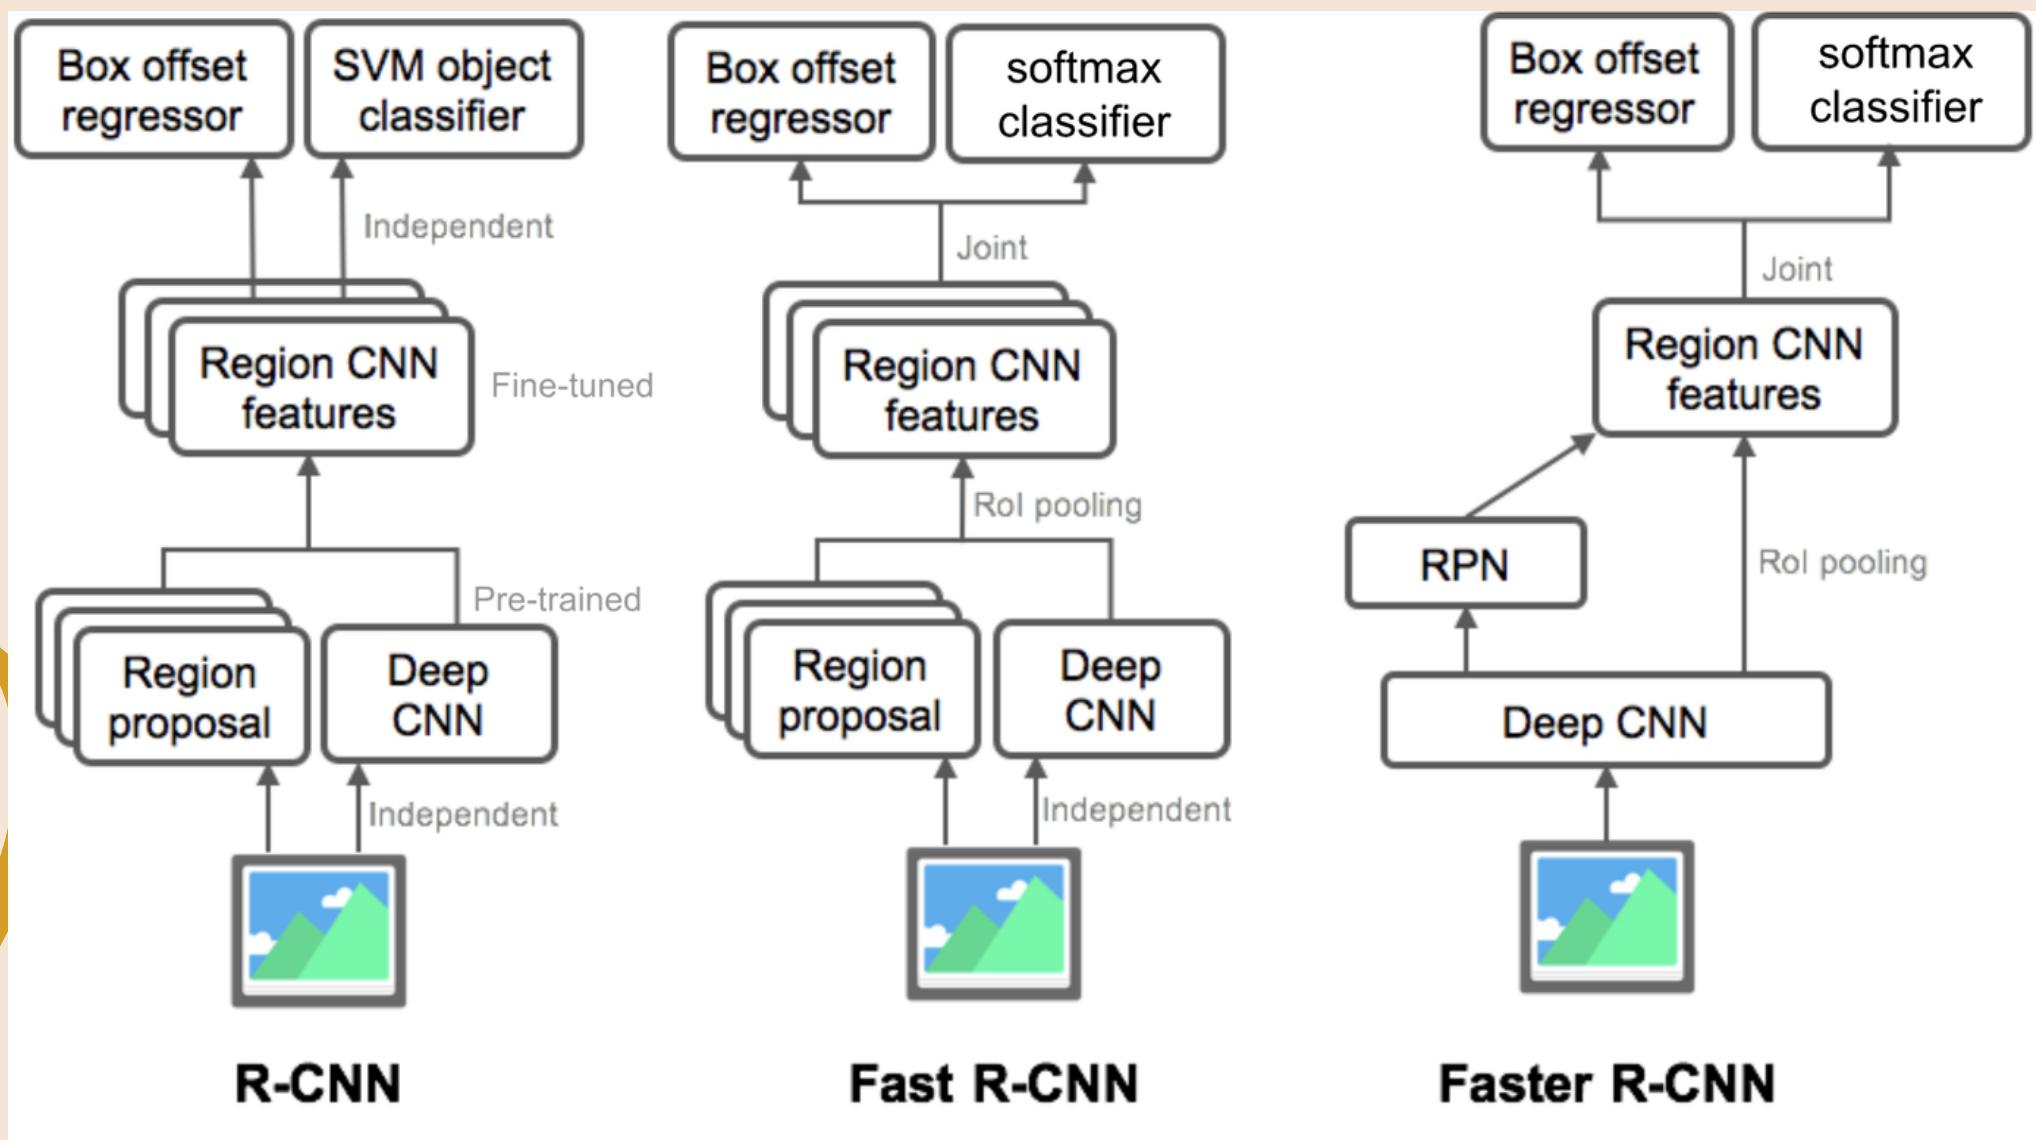
-basis: RCNN



# Baseline - image to skeleton

-optimized: Faster RCNN

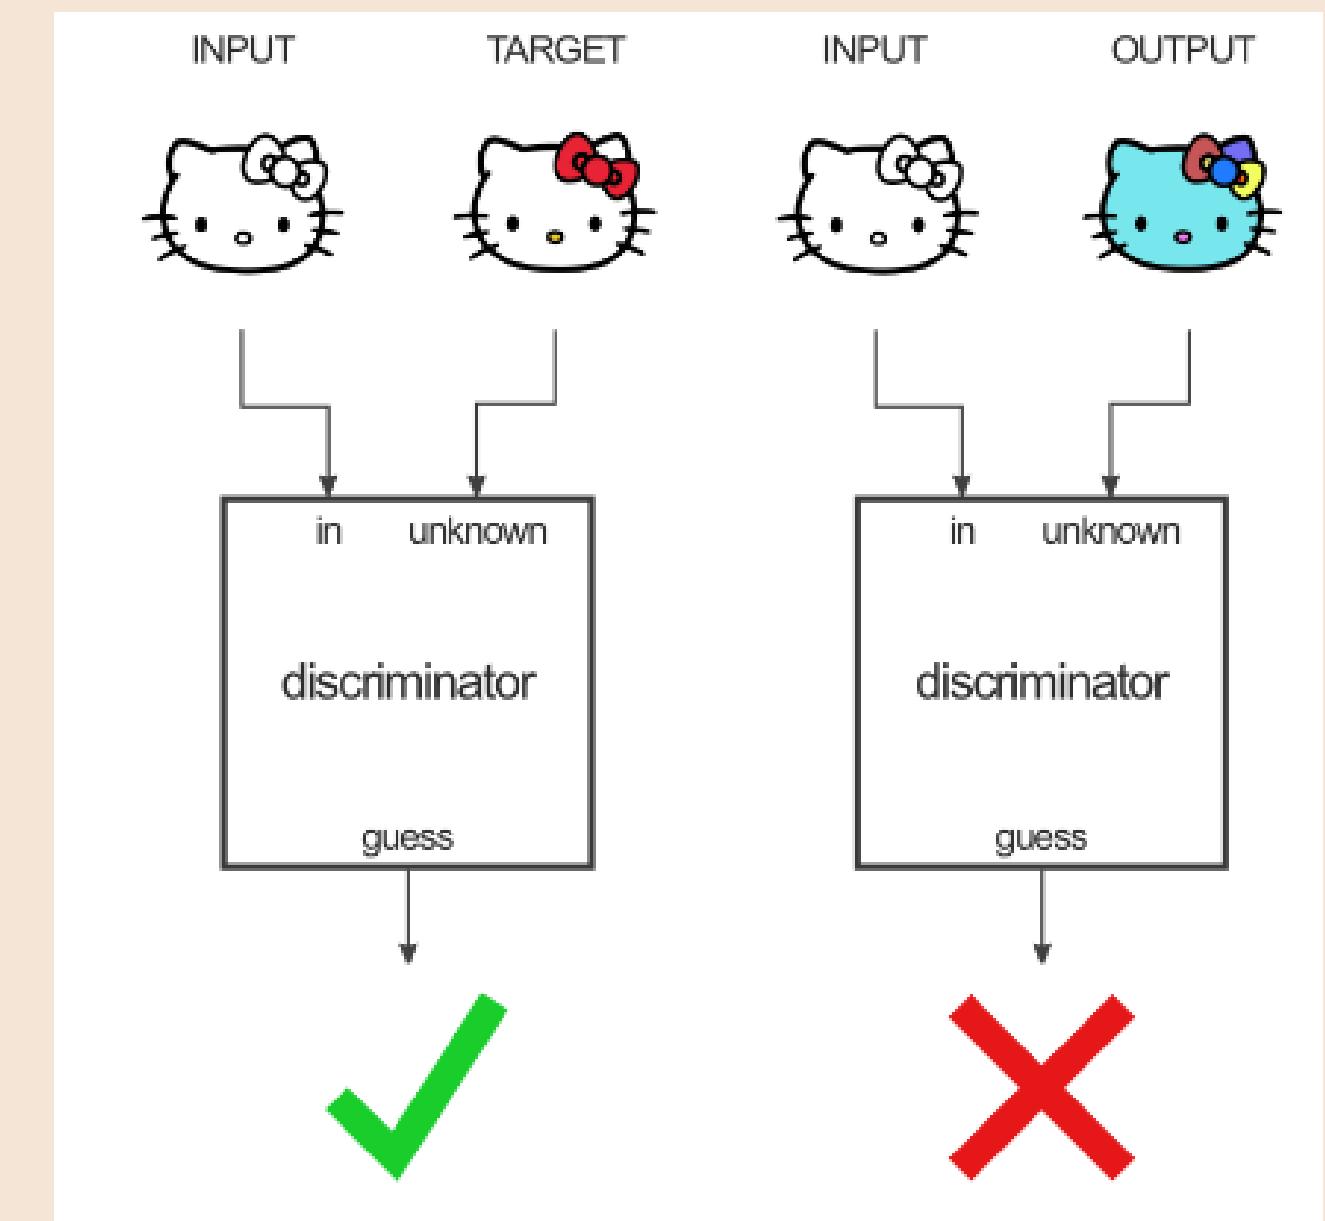
choose the most possible anchor



# Baseline - image to image

## -Conditional Adversarial Network

- generator produces image
- discriminator identifies whether the image is real or fake(generated)



# Baseline - image to image

## Conditional Adversarial Networks

generator

```
self.inc = Inconv(input_nc, ngf, norm_layer, use_bias) # input convolutional layer(64 feature maps)
self.down1 = Down(ngf, ngf * 2, norm_layer, use_bias)
self.down2 = Down([ngf * 2, ngf * 4, norm_layer, use_bias])

model = []
# 在model內建立n個殘差塊(resblock)
for i in range(n_blocks):
    model += [ResBlock(ngf * 4, padding_type=padding_type, norm_layer=norm_layer, use_dropout=use_dropout)]
self.resblocks = nn.Sequential(*model)

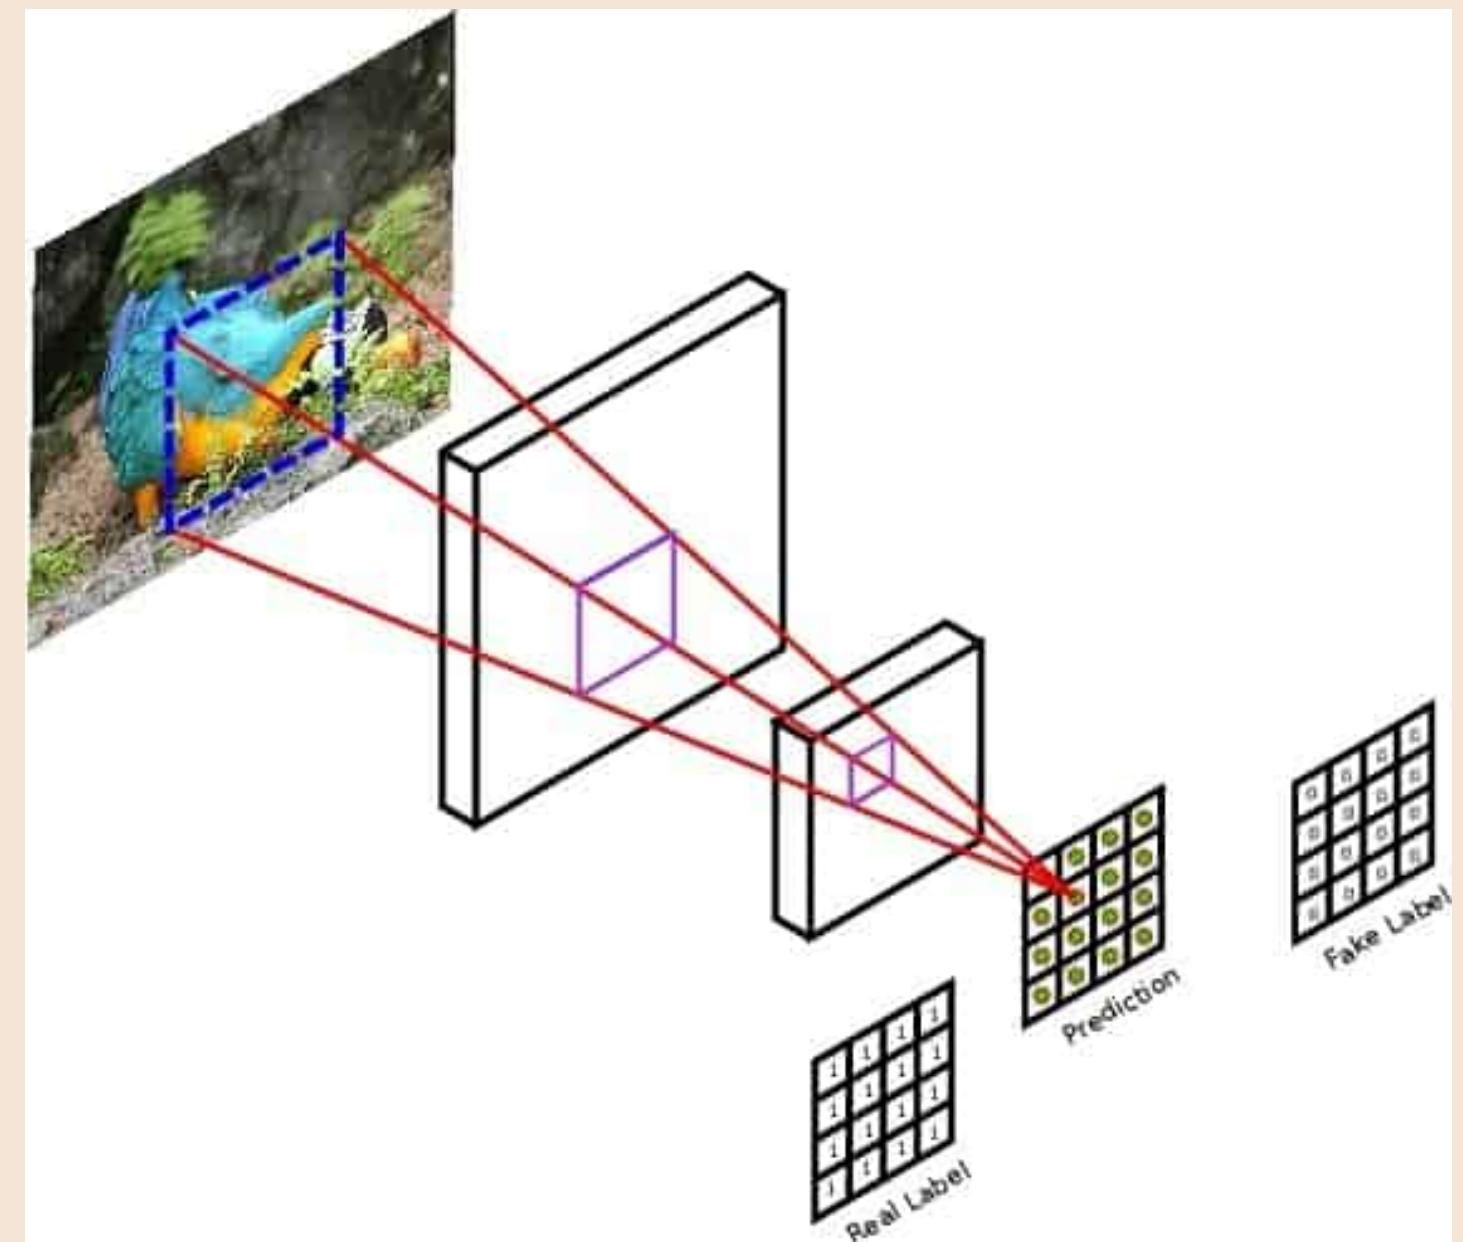
self.up1 = Up(ngf * 4, ngf * 2, norm_layer, use_bias)
self.up2 = Up(ngf * 2, ngf, norm_layer, use_bias)

self.outc = Outconv(ngf, output_nc)
```

# Baseline - image to image

-discriminator: patchGAN

- like a convolutional network
- chop the image to patches and process them independently



# Baseline - image to image

## Conditional Adversarial Networks

### discriminator

```
sequence = [
    # input channel, output channel(number of feature maps), kernel size, stride, padding(補0)
    nn.Conv2d(input_nc, ndf, kernel_size=kw, stride=2, padding=padw),
    # 上一層神經元給下一層非線性函數轉換 (助於學習)
    nn.LeakyReLU(0.2, True)
]

nf_mult = 1
nf_mult_prev = 1
for n in range(1, n_layers):
    nf_mult_prev = nf_mult
    nf_mult = min(2**n, 8)
    sequence += [
        nn.Conv2d(ndf * nf_mult_prev, ndf * nf_mult,
                 kernel_size=kw, stride=2, padding=padw, bias=use_bias),
        norm_layer(ndf * nf_mult),
        nn.LeakyReLU(0.2, True)
    ]
```

# MAIN APPROACH

## Keypoint Detection

- Model - based on FasterRCNN
- Train - with some optimizers
- Visualize the result

## Image to Image

- Model - GAN
- Train - with some optimizers
- Generate image

# MAIN APPROACH

## Keypoint Detection

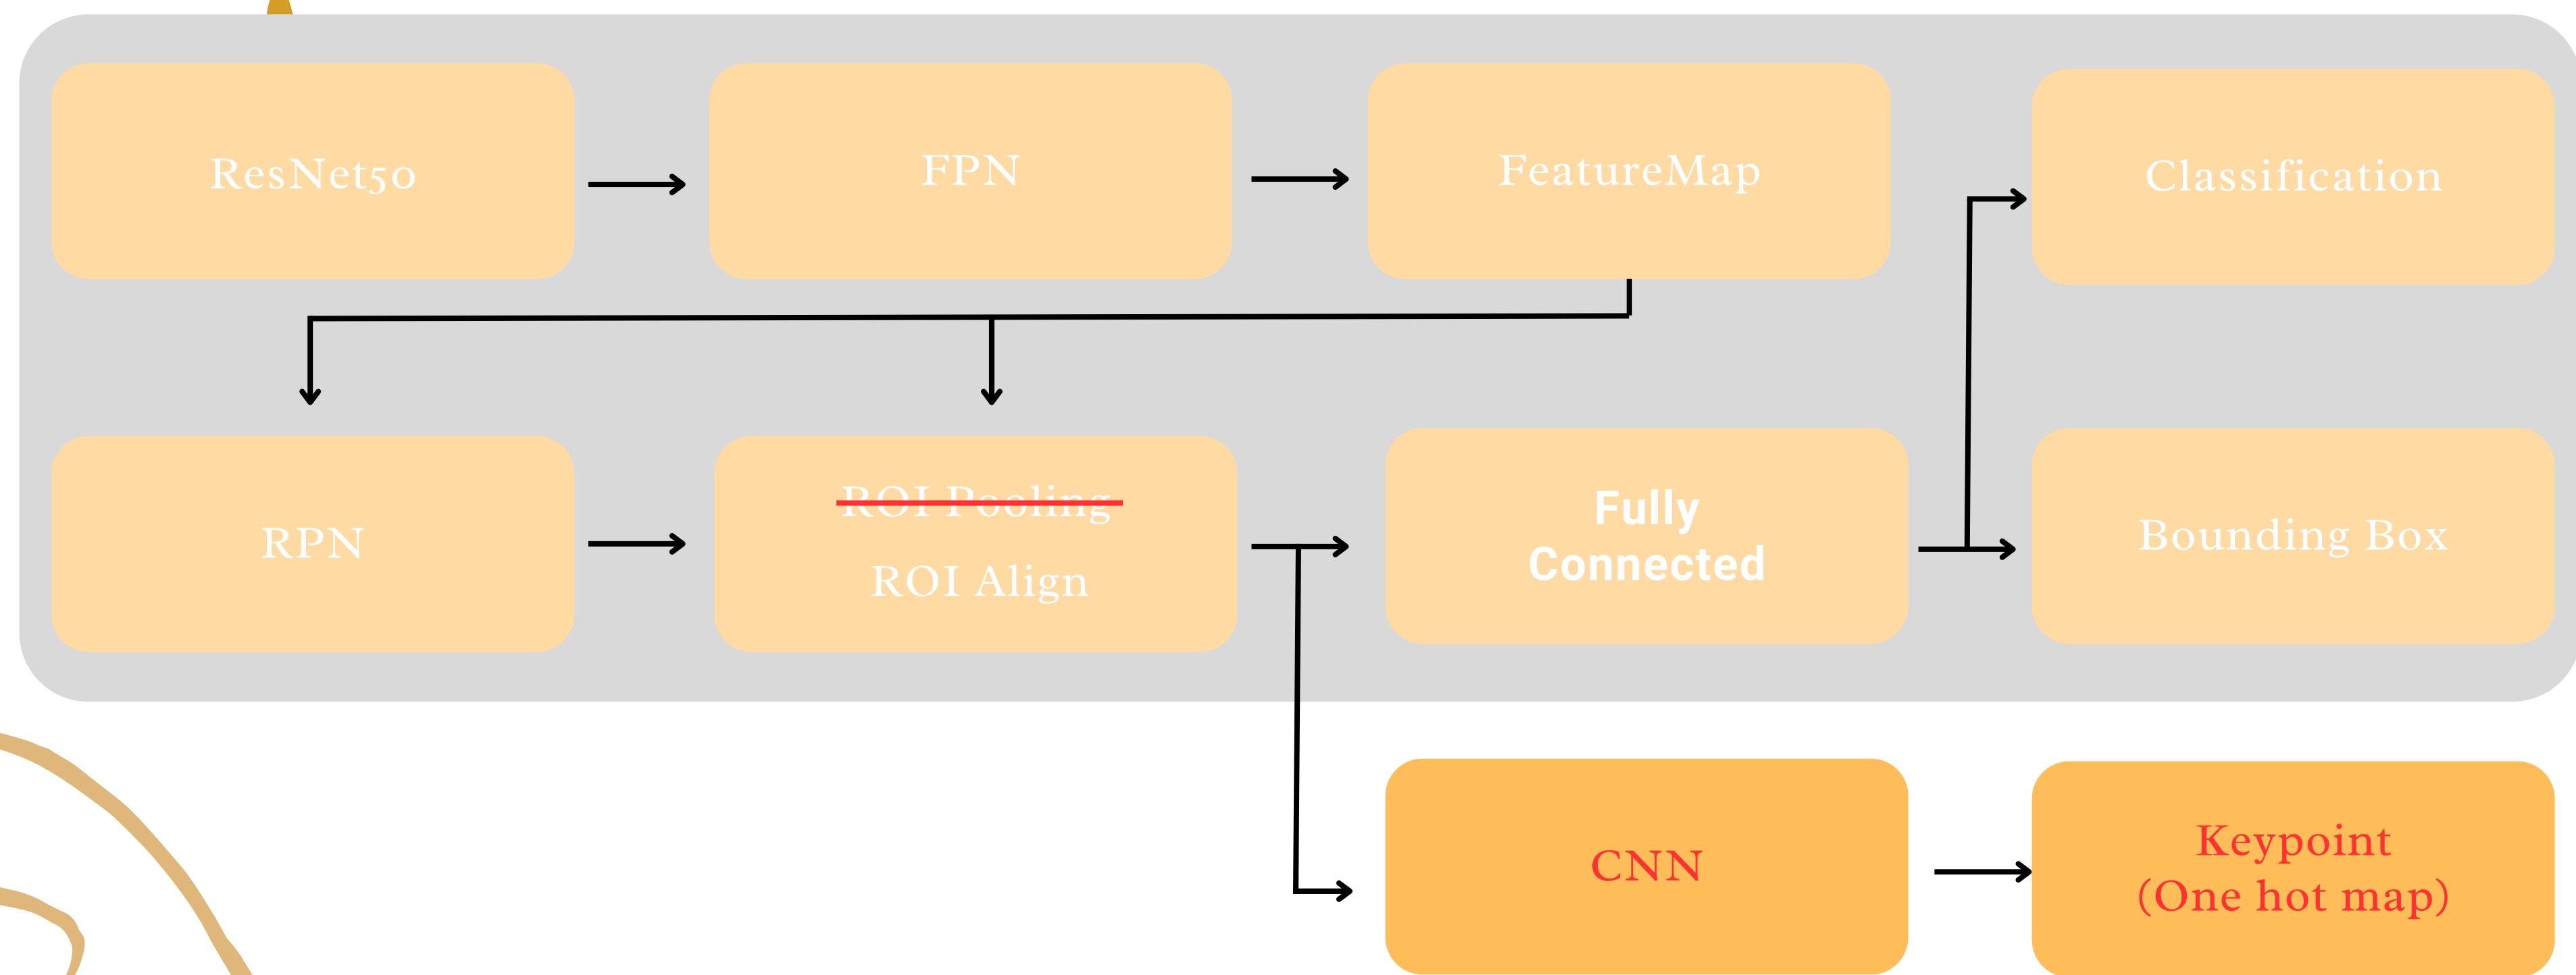
- Model - based on FasterRCNN
- Train - with some optimizers
- Visualize the results

## Image to Image

- Model - GAN
- Train - with some optimizers
- Generate image

# Model

Faster RCNN



# Model-CNN

```
def __init__(self):
    super(KeypointRCNNPredictor, self).__init__()
    self.kps_score_lowres = nn.ConvTranspose2d(512, 17, kernel_size=4, stride=2, padding=1)
```

Batch Normalization: normalizes the inputs of each mini-batch

- accelerate convergence
  - keep inputs distribution stable at each layer
  - improve stability and accuracy

`ConvTranspose2d`: increases the size of a feature map

- convert low-resolution feature maps to high-resolution ones
  - aiding in detail restoration and improving spatial accuracy.

# Train

## STEP I: Loading data

- Transform: increasing the diversity of training data
- Using RandomSampler to randomly samples elements from the dataset
- Using BatchSampler to divide the samples into batches then using data loader
- Finally we can use the dataset from the data loader to train the model

```
def get_transform(train):
    transforms = []
    transforms.append(T.ToTensor())
    if train:
        transforms.append(T.RandomHorizontalFlip(0.5))
    return T.Compose(transforms)
```

```
num = list(range(1, len(dataset_test), 2))
dataset_test = torch.utils.data.Subset(dataset_test, num)

train_sampler = torch.utils.data.RandomSampler(dataset)
test_sampler = torch.utils.data.RandomSampler(dataset_test)
train_batch_sampler = torch.utils.data.BatchSampler(train_sampler, args.batch_size, drop_last=True)
```

# Train

## STEP2: Optimize and train the model

- Using SGD + momentum: the update step for the model parameters is determined by both the current gradient and the accumulated momentum from previous iterations.
- learning rate scheduler: multiplies the learning rate by a specified factor at predetermined steps ->  
 $lr = lr * \text{gamma}$
- Training the model

```
# optimize the model
params = [p for p in model.parameters() if p.requires_grad]
optimizer = torch.optim.SGD(params, lr=args.lr, momentum=args.momentum, weight_decay=args.weight_decay)

#adjust learning rate every step_size lr = lr*gamma
lr_scheduler = torch.optim.lr_scheduler.StepLR(optimizer, step_size=args.lr_step_size, gamma=args.lr_gamma)
```

```
for epoch in range(args.epochs):
    train_one_epoch(model, optimizer, data_loader, device, epoch, args.print_freq)
    lr_scheduler.step()
    evaluate(model, data_loader_test, device=device) # evaluate after every epoch
    torch.save(model.state_dict(), 'training_weights'+str(i)+'.pth') # save model /epoch
```

# Visualize

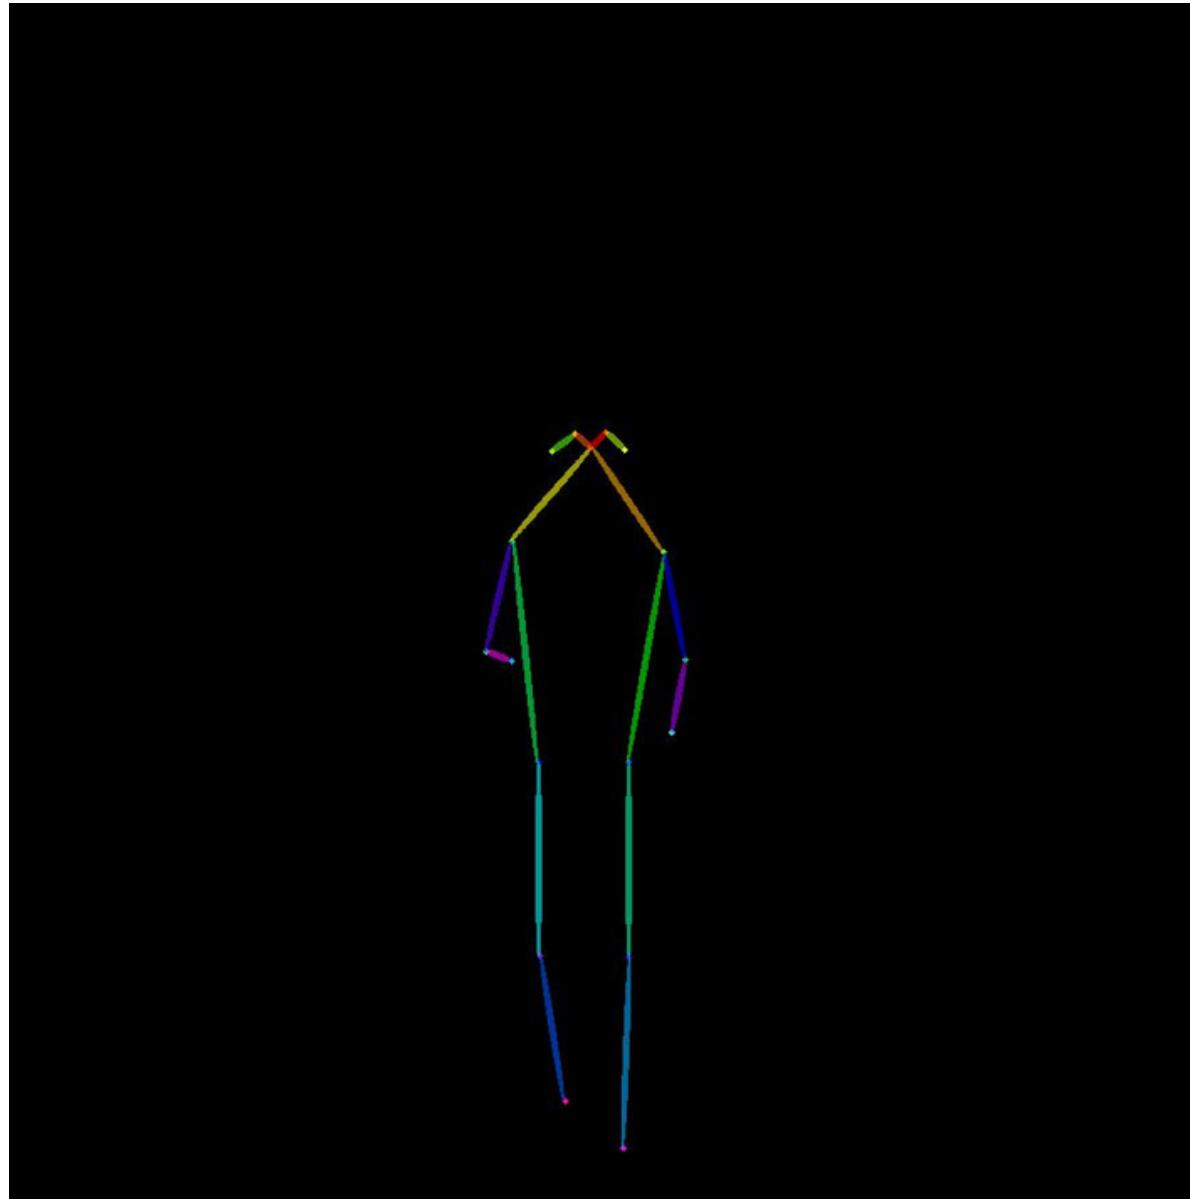
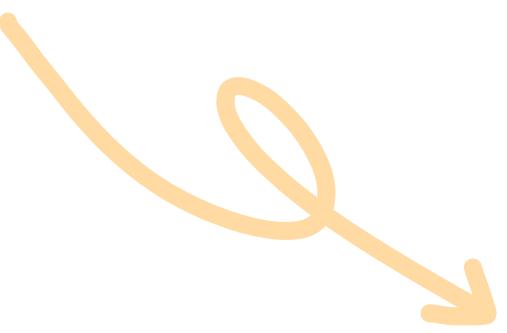
## STEP3: Plot the results

- Load the training weight for model
- Using detect\_threshold, keypoint\_score\_threshold parameters to filter out individuals with low scores and keypoints with low scores,

```
model = get_model(train_weight='training_weights.pth')
```

```
detect_threshold = 0.7
keypoint_score_threshold = 2
with torch.no_grad():
    for i in range(10):
        print(i)
        img, _ = dataset_test[i]
        prediction = model([img.to(device)])
        keypoints = prediction[0]['keypoints'].cpu().numpy()
        scores = prediction[0]['scores'].cpu().numpy()
        keypoints_scores = prediction[0]['keypoints_scores'].cpu().numpy()
        idx = np.where(scores>detect_threshold)
        keypoints = keypoints[idx]
        keypoints_scores = keypoints_scores[idx]
        for j in range(keypoints.shape[0]):
            for num in range(17):
                if keypoints_scores[j][num]<keypoint_score_threshold:
                    keypoints[j][num]=[0,0,0]
        img = img.mul(255).permute(1, 2, 0).byte().numpy()
        plot_poses(img, keypoints, save_name='./result/'+str(i)+'.jpg')
print('Finish!')
```

# Visualize



# MAIN APPROACH

## Keypoint Detection

- Model - based on FasterRCNN
- Train - with some optimizers
- Visualize the result

## Image to Image

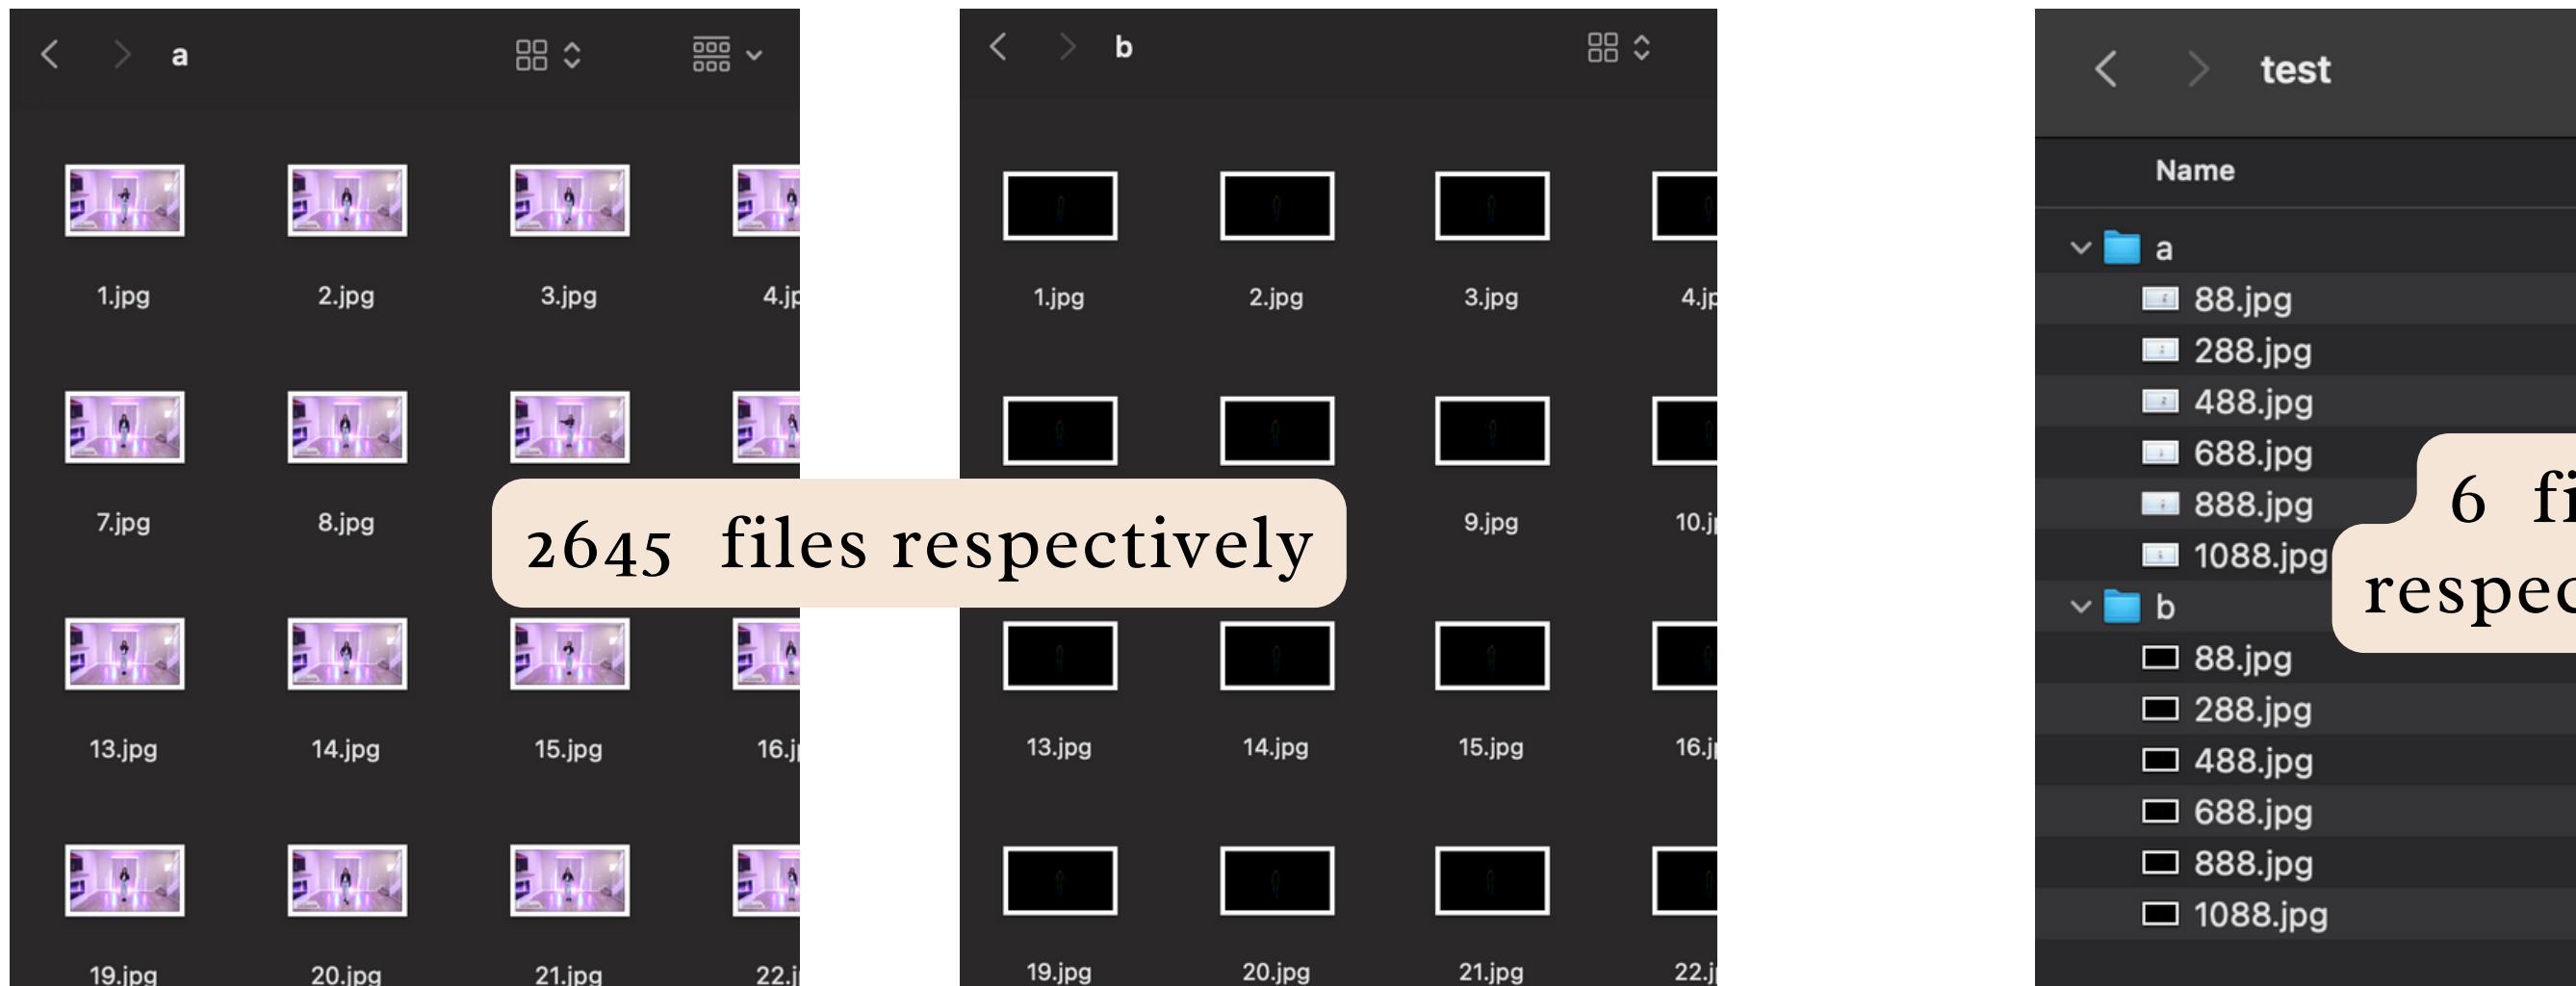
- Model - GAN
- Train - with some optimizers
- Generate image

# Train GAN model

## STEP I : Loading Dataset

- training data : 2645 files, testing data : 6 files

```
print('====> Loading datasets')
root_path = "dataset1/"
train_set = get_training_set(root_path + opt.dataset, opt.direction)
test_set = get_test_set(root_path + opt.dataset, opt.direction)
training_data_loader = DataLoader(dataset=train_set, num_workers=opt.threads, batch_size=opt.batch_size, shuffle=True)
testing_data_loader = DataLoader(dataset=test_set, num_workers=opt.threads, batch_size=opt.test_batch_size, shuffle=False)
```



# Train GAN model

## STEP 2 : Model Definition

- Define Generator, Discriminator, GAN loss function

```
print('====> Building models')
net_g = define_G(opt.input_nc, opt.output_nc, opt.ngf, 'batch', False, 'normal', 0.02, gpu_id=device)
net_d = define_D(opt.input_nc + opt.output_nc, opt.ndf, 'basic', gpu_id=device)

criterionGAN = GANLoss().to(device)
criterionL1 = nn.L1Loss().to(device)
criterionMSE = nn.MSELoss().to(device)
```

- Setup Optimizer and learning rate scheduler

```
optimizer_g = optim.Adam(net_g.parameters(), lr=opt.lr, betas=(opt.beta1, 0.999))
optimizer_d = optim.Adam(net_d.parameters(), lr=opt.lr, betas=(opt.beta1, 0.999))
net_g_scheduler = get_scheduler(optimizer_g, opt)
net_d_scheduler = get_scheduler(optimizer_d, opt)
```

# Train GAN model

## Introduction to optimizer

SGD  
(stochastic gradient decent)

Momentum SGD

Adam

calculate weight : gradient of the loss function by differentiation  
update parameters : opposite direction of gradient and learning rate

$$W \leftarrow W - \eta \frac{\partial L}{\partial W}$$

allowing for faster updates in the same direction

$$V_t \leftarrow \beta V_{t-1} - \eta \frac{\partial L}{\partial W}$$

Keeps the exponentially decaying average of the past gradient  
keeps the squared decaying average of the past gradient

# Train GAN model

## STEP3 : Training Loop

- Updates the discriminator

```
for iteration, batch in enumerate(training_data_loader, 1): # forward
    real_a, real_b = batch[0].to(device), batch[1].to(device)
    fake_b = net_g(real_a)

    #####
    # (1) Update D network
    #####
    optimizer_d.zero_grad()

    # train with fake
    fake_ab = torch.cat((real_a, fake_b), 1)
    pred_fake = net_d.forward(fake_ab.detach())
    loss_d_fake = criterionGAN(pred_fake, False)

    # train with real
    real_ab = torch.cat((real_a, real_b), 1)
    pred_real = net_d.forward(real_ab)
    loss_d_real = criterionGAN(pred_real, True)

    # Combined D loss
    loss_d = (loss_d_fake + loss_d_real) * 0.5
    loss_d.backward()

    optimizer_d.step()
```

## generator

- loss calculations are performed for the generated fake images and real images and the model weights are updated based on these losses.

```
#####
# (2) Update G network
#####

optimizer_g.zero_grad()

# First, G(A) should fake the discriminator
fake_ab = torch.cat((real_a, fake_b), 1)
pred_fake = net_d.forward(fake_ab)
loss_g_gan = criterionGAN(pred_fake, True)

# Second, G(A) = B
loss_g_l1 = criterionL1(fake_b, real_b) * opt.lamb

loss_g = loss_g_gan + loss_g_l1
loss_g.backward()

optimizer_g.step()

print("====> Epoch[{}]({})/{}: Loss_D: {:.4f} Loss_G: {:.4f}".format(
    epoch, iteration, len(training_data_loader), loss_d.item(), loss_g.item()))
```

# Train GAN model

## STEP4 : Model preservation and testing

| Filename                                | Filesize | Filetype |
|---|----------|----------|
| ..                                      |          |          |
| netG_model_epoch_20.pth                 | 45637225 | pth-file |
| netD_model_epoch_20.pth                 | 11091427 | pth-file |
| netG_model_epoch_10.pth                 | 45637225 | pth-file |
| netD_model_epoch_10.pth                 | 11091427 | pth-file |
| netG_model_epoch_200.pth                | 45637225 | pth-file |
| netD_model_epoch_200.pth                | 11091427 | pth-file |
| netG_model_epoch_190.pth                | 45637225 | pth-file |
| netD_model_epoch_190.pth                | 11091427 | pth-file |
| netG_model_epoch_180.pth                | 45637225 | pth-file |
| netD_model_epoch_180.pth                | 11091427 | pth-file |
| netG_model_epoch_170.pth                | 45637225 | pth-file |
| netD_model_epoch_170.pth                | 11091427 | pth-file |
| netG_model_epoch_160.pth                | 45637225 | pth-file |
| netD_model_epoch_160.pth                | 11091427 | pth-file |
| netG_model_epoch_150.pth                | 45637225 | pth-file |
| netD_model_epoch_150.pth                | 11091427 | pth-file |
| netG_model_epoch_140.pth                | 45637225 | pth-file |
| netD_model_epoch_140.pth                | 11091427 | pth-file |
| netG_model_epoch_130.pth                | 45637225 | pth-file |
| netD_model_epoch_130.pth                | 11091427 | pth-file |
| netG_model_epoch_120.pth                | 45637225 | pth-file |
| 120 files. Total size: 3403719120 bytes |          |          |

```
# test
avg_psnr = 0
for batch in testing_data_loader:
    input, target = batch[0].to(device), batch[1].to(device)

    prediction = net_g(input)
    mse = criterionMSE(prediction, target)
    psnr = 10 * log10(1 / mse.item())
    avg_psnr += psnr
print("====> Avg. PSNR: {:.4f} dB".format(avg_psnr / len(testing_data_loader)))
```

Performance metric : PSNR (peak signal-to-noise ratio)  
a common measure of image quality.

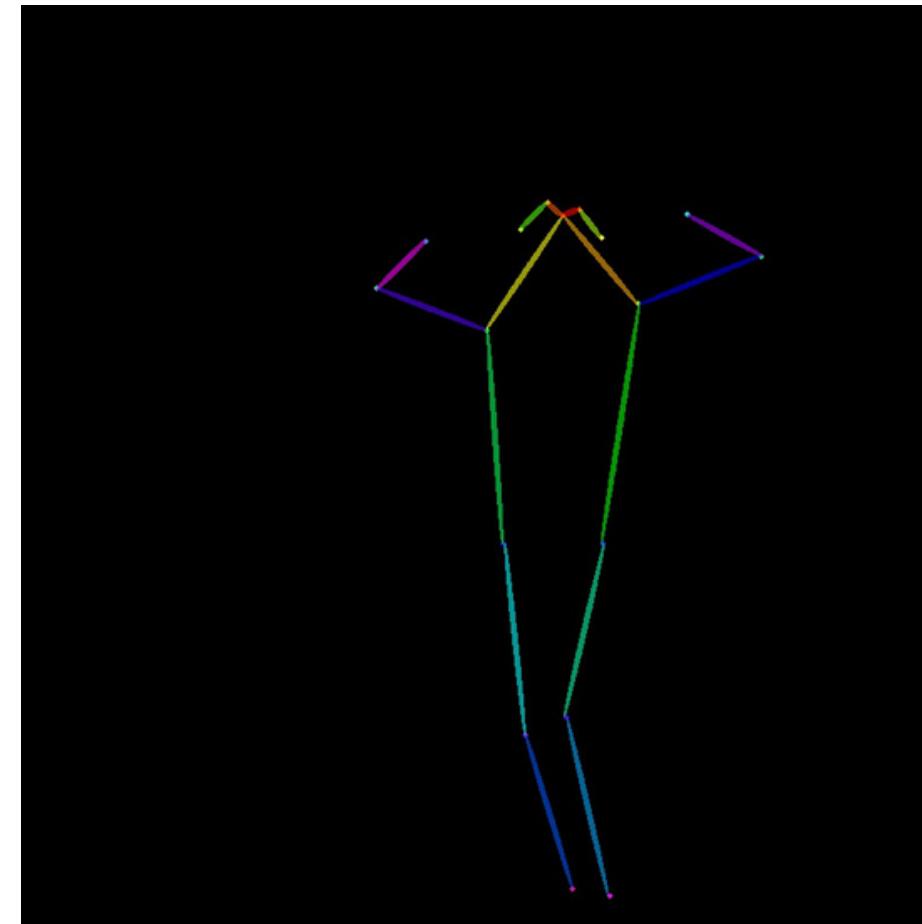
```
====> Epoch[15](2643/2645): Loss_D: 0.0014 Loss_G: 2.1399
====> Epoch[15](2644/2645): Loss_D: 0.0487 Loss_G: 2.1273
====> Epoch[15](2645/2645): Loss_D: 0.0292 Loss_G: 2.2286
learning rate = 0.0008000
learning rate = 0.0008000
====> Avg. PSNR: 6.9454 dB
====> Epoch[16](1/2645): Loss_D: 0.0153 Loss_G: 2.0304
====> Epoch[16](2/2645): Loss_D: 0.0111 Loss_G: 1.9949
====> Epoch[16](3/2645): Loss_D: 0.0031 Loss_G: 1.9677
```

# Image Generation

STEP : Image Transform



Test Dataset



Result



# Evaluation Metric

## DATA ACCURACY

- IoU = 0.5
- Precision
- Compare the accuracy at each epoch

## RUNNING TIME

- End time -Start time

# Results & Analysis

## Keypoint Detection:

- Learning rate
- Optimizer: SGD, SGD+momentum, Adagrade
- add BatchNorm2d

## Image to Image

- find best optimizer
- find optimal learning rate
- find optimal epoch

# Results & Analysis

## Keypoint Detection:

- Learning rate
- Optimizer: SGD, SGD+momentum, Adagrade
- add BatchNorm2d

## Image to Image

- find best optimizer
- find optimal learning rate
- find optimal epoch

# Results & Analysis

## OPTIMAL LEARNING RATE

### Result

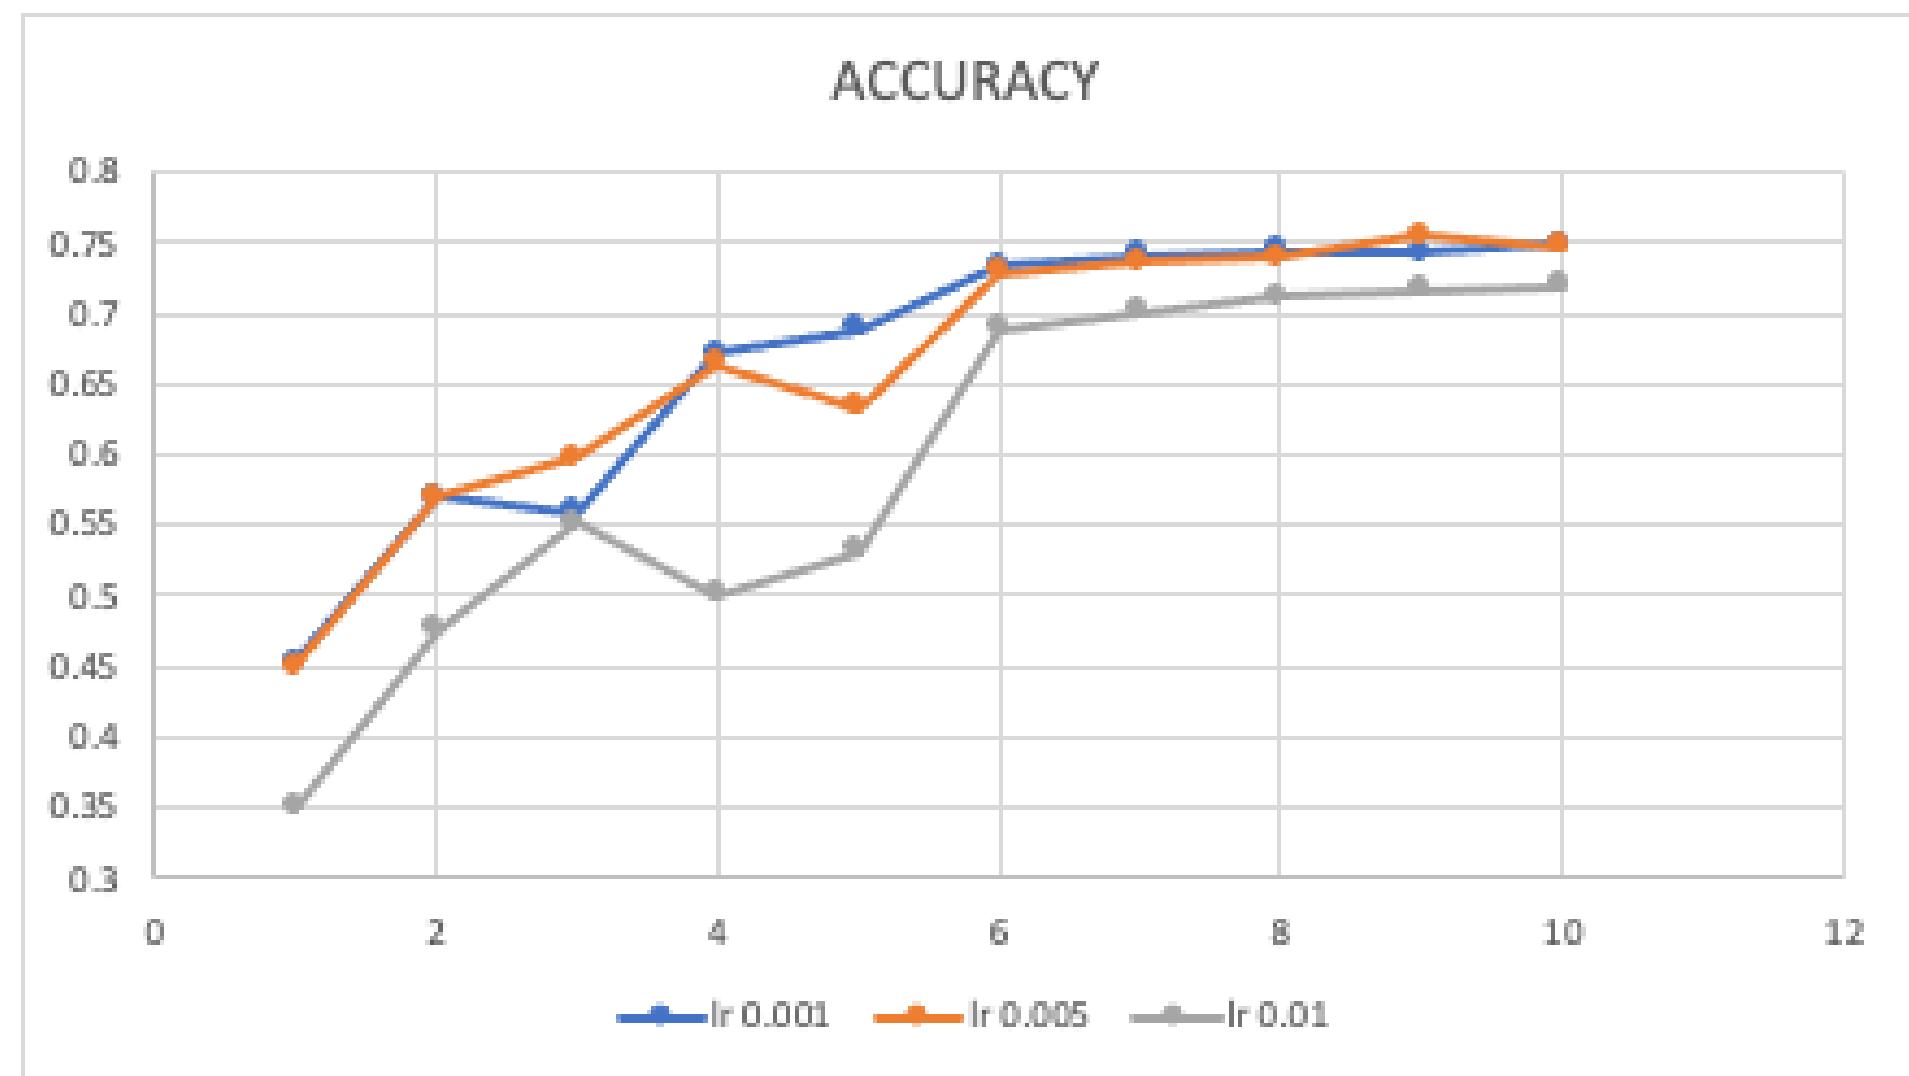
Best performance:

learning rate = 0.001, 0.005

### Analysis

With learning rates 0.001 and 0.005, the accuracies are relatively high and show less fluctuation across iterations.

In contrast, the accuracy is lower for the learning rate of 0.01, and there is more significant fluctuation across iterations.



# Results & Analysis

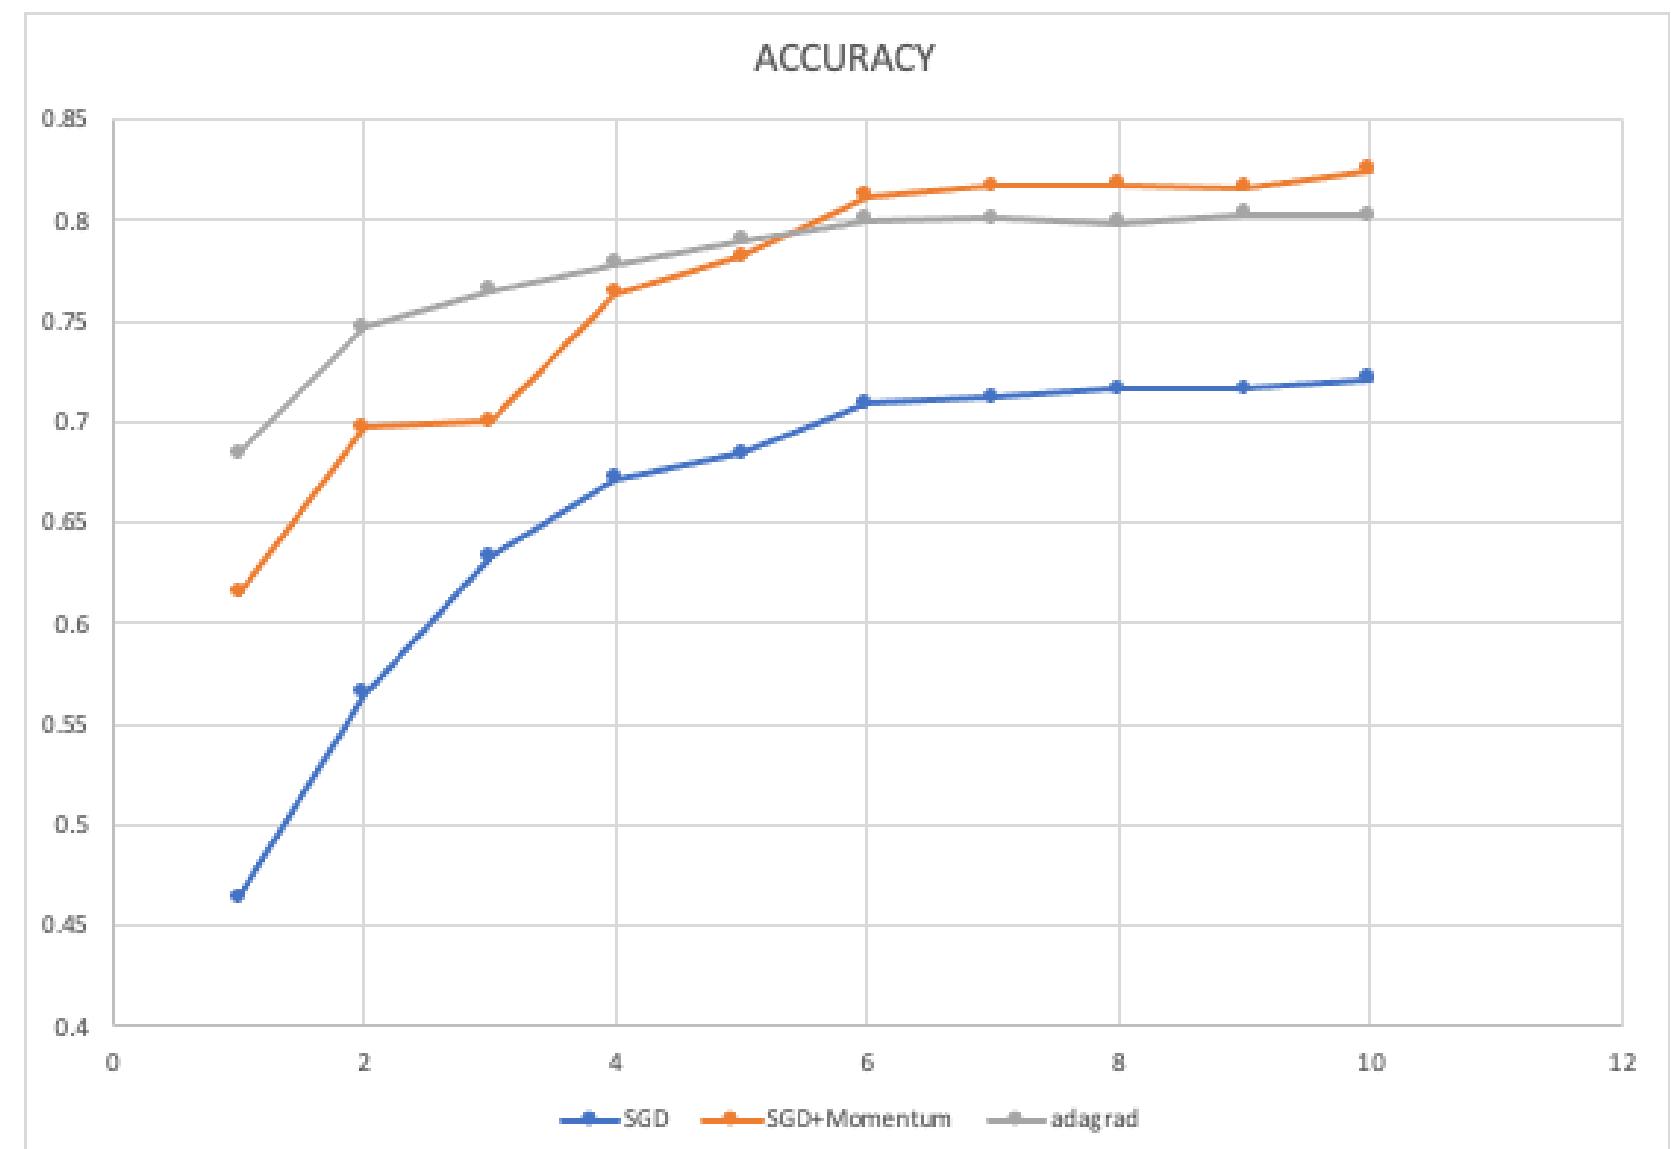
## OPTIMIZER: Adagrad v.s SGD v.s SGD+Momentum

### Result

Best performance:  
SGD+Momentum

### Analysis

- Adagrad's characteristic is to adaptively adjust the learning rate based on the historical gradient information of parameters. It may cause the learning rate to decrease excessively
- SGD+Momentum optimizer incorporates the concept of momentum, which helps it better handle the direction and magnitude of gradients during optimization



# Results & Analysis

## OPTIMIZER: Adagrad v.s SGD v.s SGD+Momentum

### Result

SGD+Momentum has the shortest training time

### Analysis

- SGD+momentum could be due to the introduction of the momentum term.
- adagrad: lead to excessive learning rate decay and slower convergence



# Results & Analysis

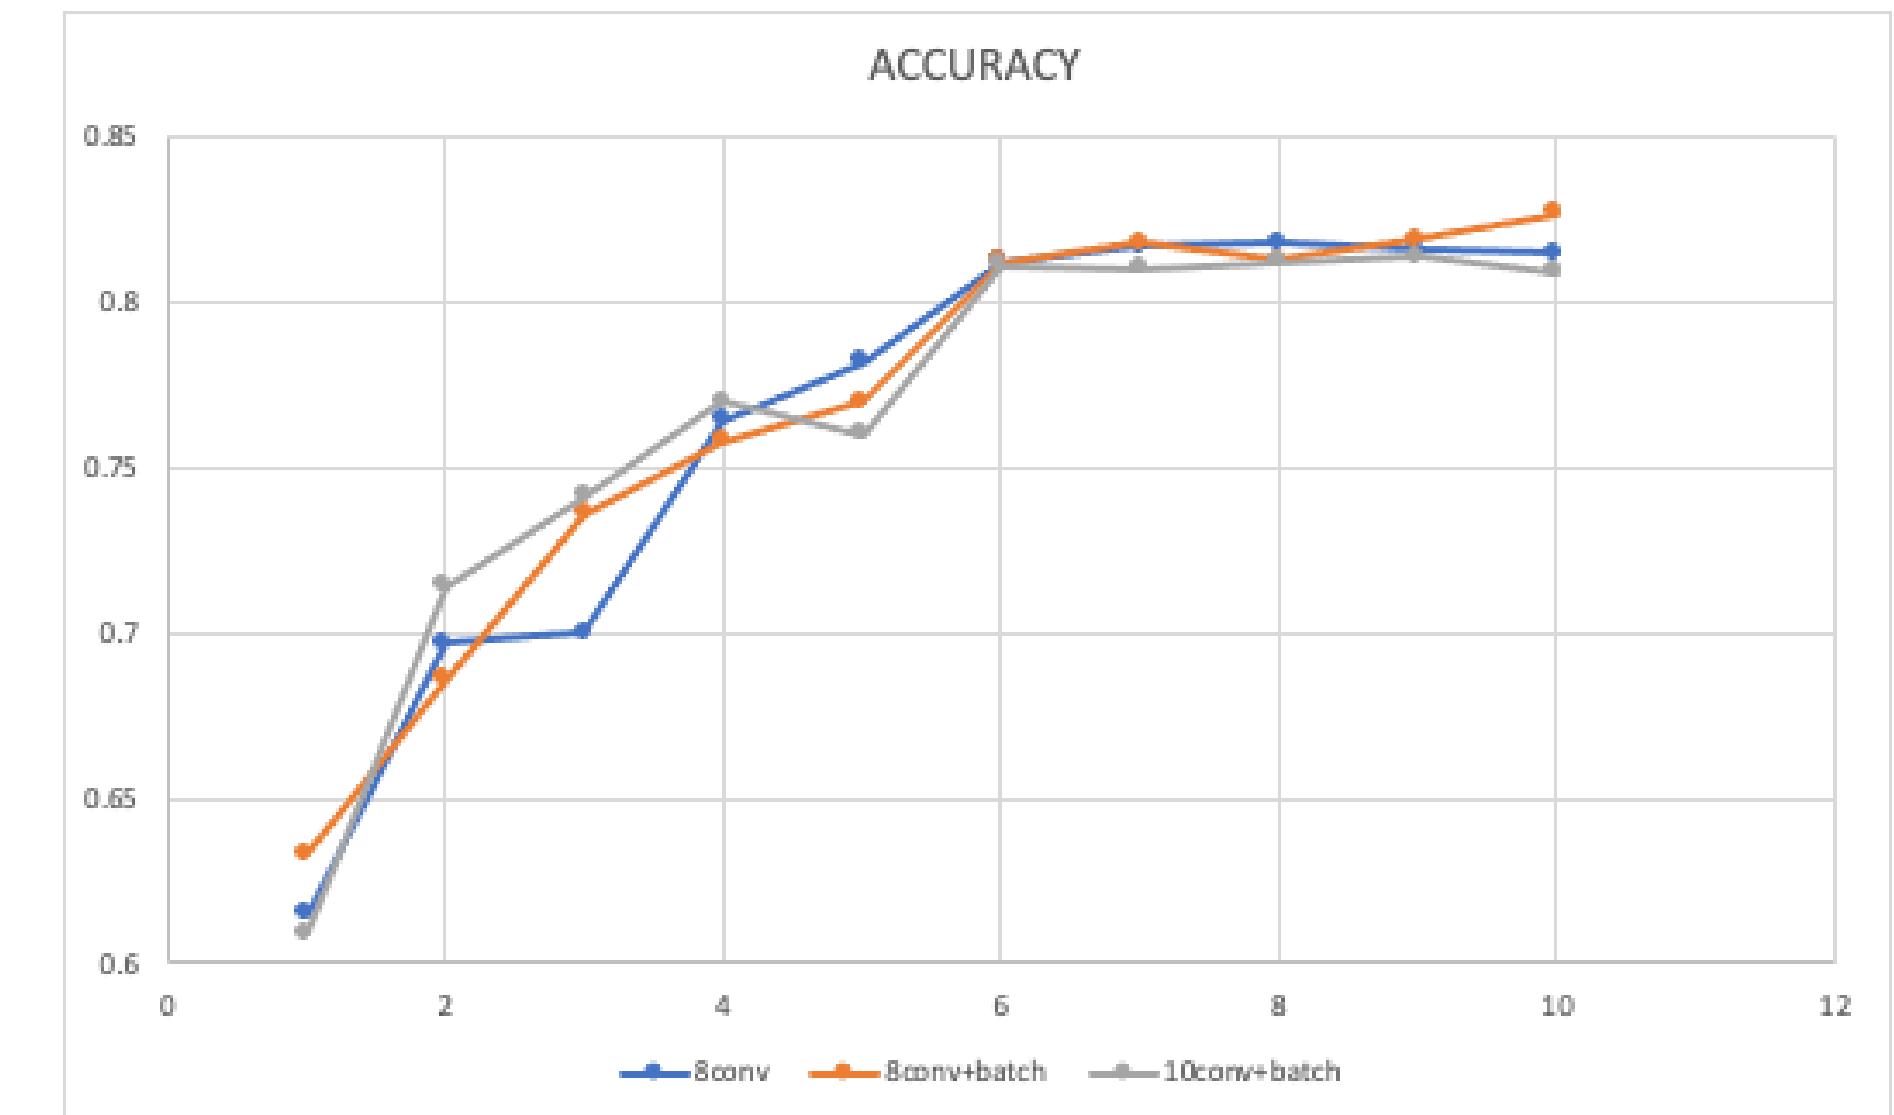
Layer: 8\*conv2d v.s 8\*conv2d + BatchNorm2d layer v.s 10\*conv2+BatchNorm2d

## Result

8\*conv2d+batchNorm2d>  
8\*conv2d>  
10\*conv2d+batchNorm2d

## Analysis

8\*conv2d+batchNorm2d: adding batch normalization improves the model's performance.



# Results & Analysis

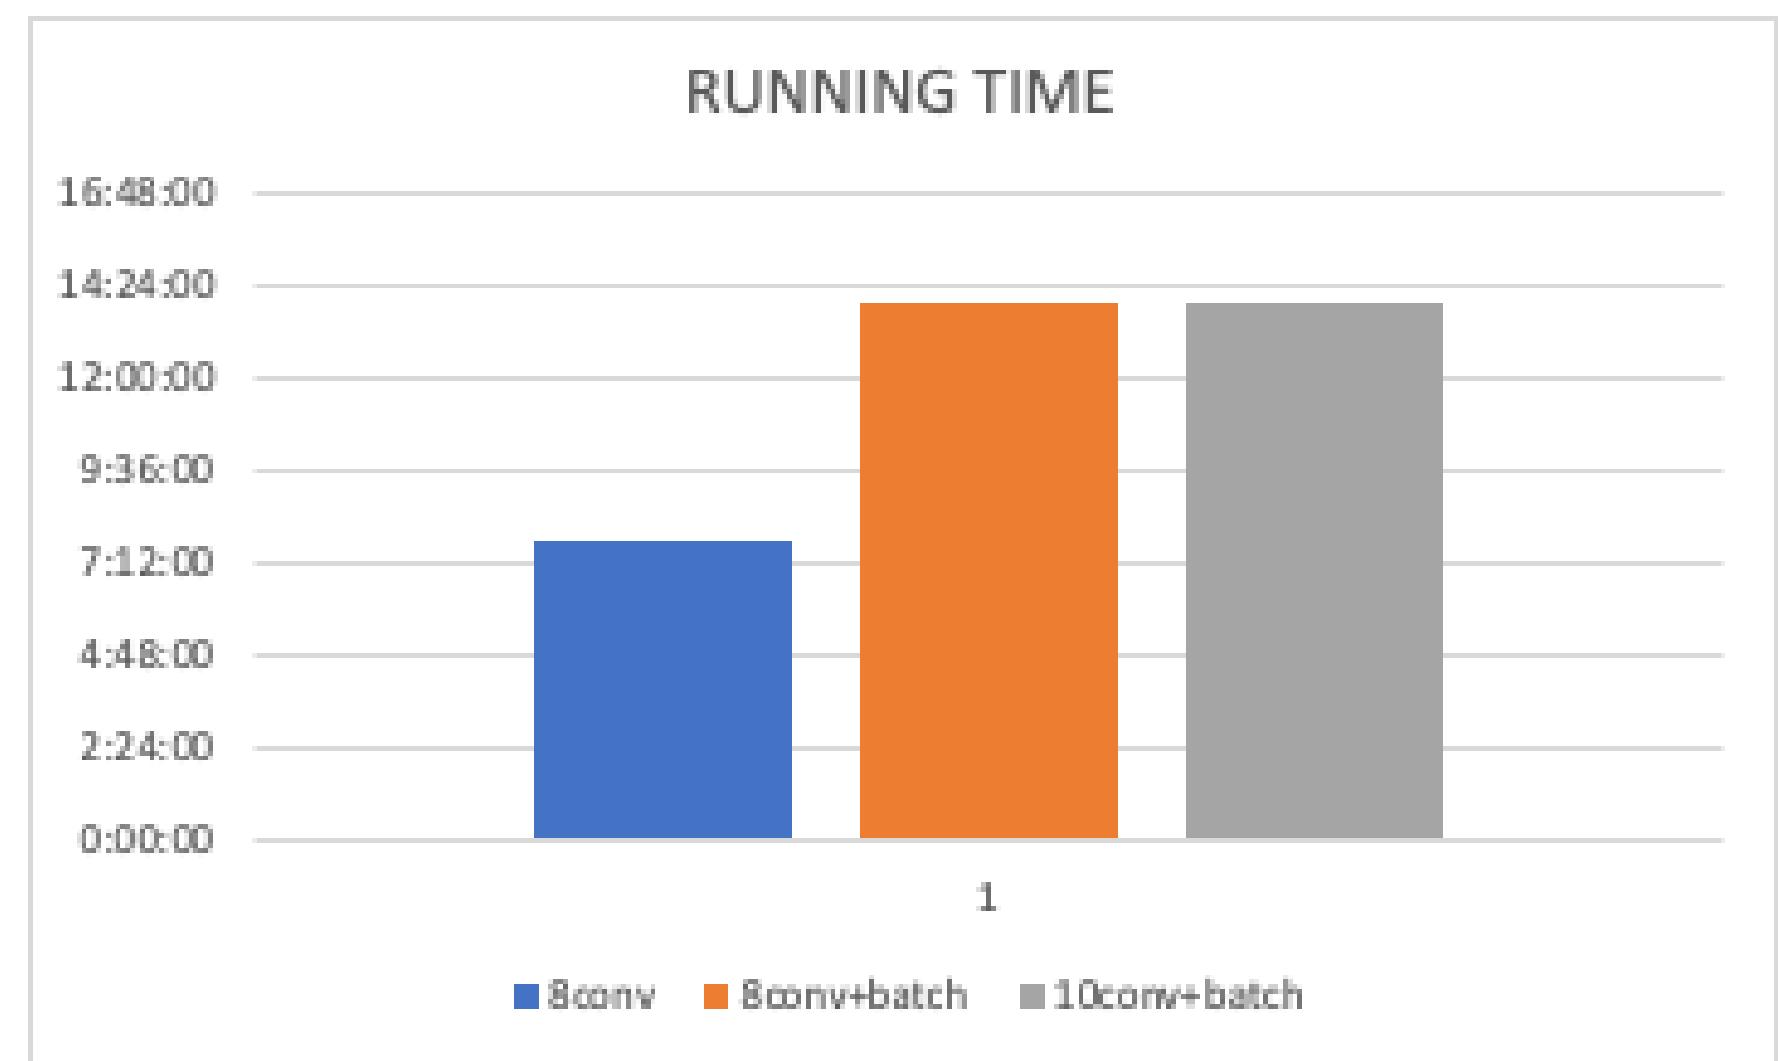
Layer: 8\*conv2d v.s 8\*conv2d + BatchNorm2d layer v.s 10\*conv2+BatchNorm2d

## Result

8\*conv2d has the shortest training time

## Analysis

- The 8conv model has a shorter training time due to its fewer convolutional layers and parameters.
- The 8conv+batch and 10conv+batch models have longer training times due to their higher number of convolutional layers and parameters.



# Results & Analysis

## Keypoint Detection:

- Learning rate
- Optimizer: SGD, SGD+momentum, Adagrade
- add BatchNorm2d

## Image to Image

- find best optimizer
- find optimal learning rate
- find optimal epoch

# Results & Analysis

Task I: Optimizer

Momentum SGD

2 Epoch



40 Epoch



Adam



test data : 488.jpg  
learning rate : 0.0002

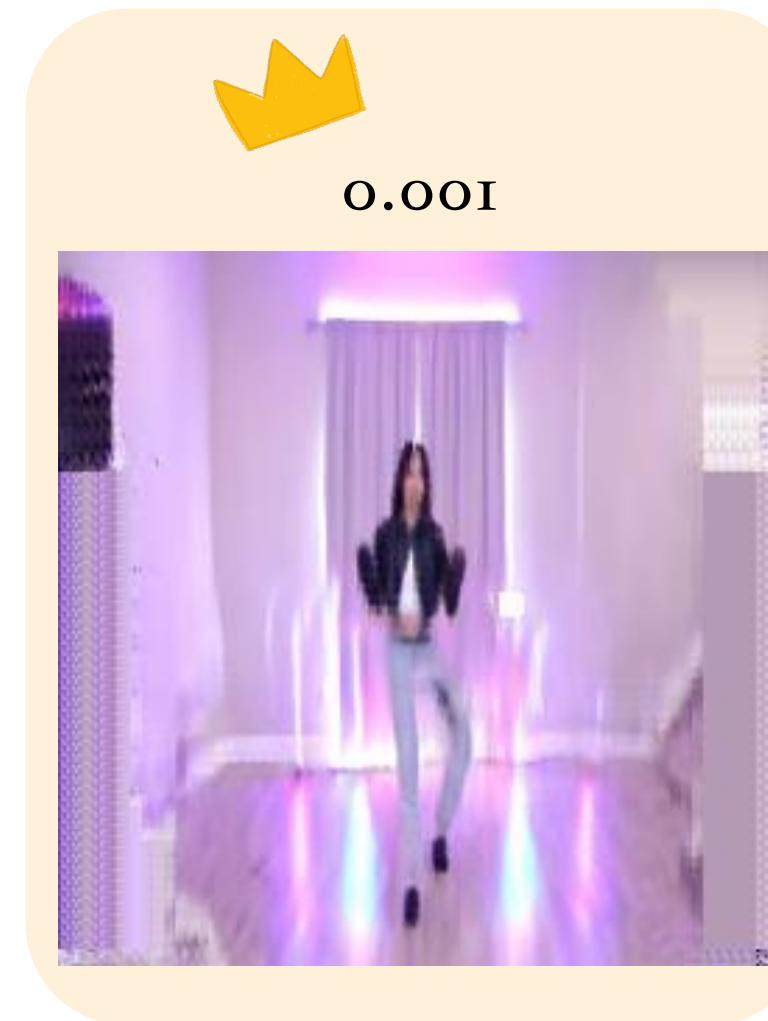
# Results & Analysis

## Task 2: Learning Rate

20 Epoch



0.0004



0.001



0.002



0.004

test data : 1088.jpg  
optimizer : Adam

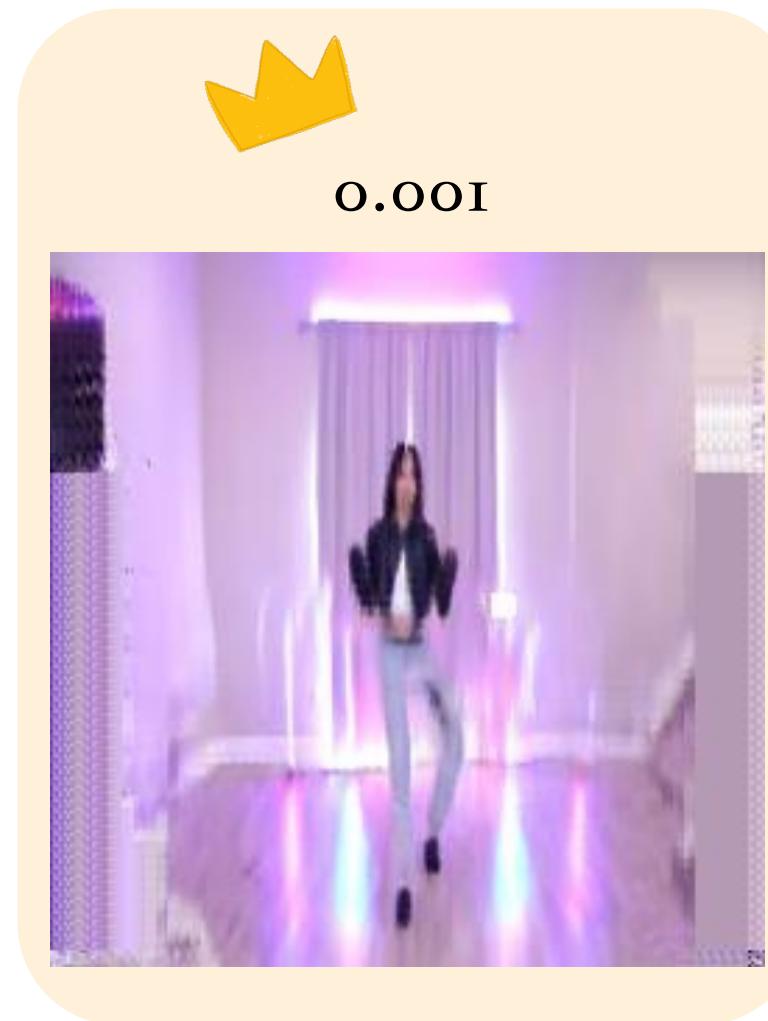
# Results & Analysis

## Task 2: Learning Rate

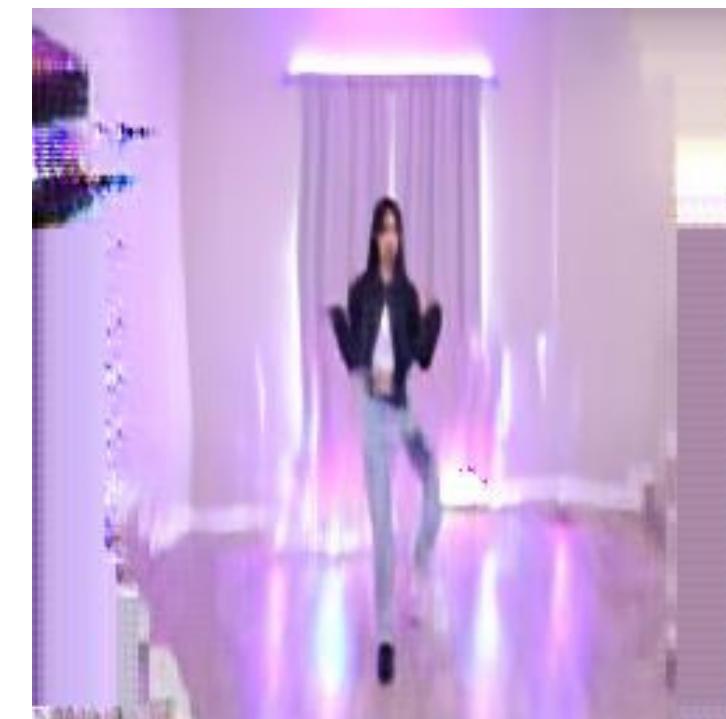
20 Epoch



0.0008



0.001



0.0012



0.0015

test data : 1088.jpg  
optimizer : Adam

# Results & Analysis

Task 3: Epoch

lr 0.001

40



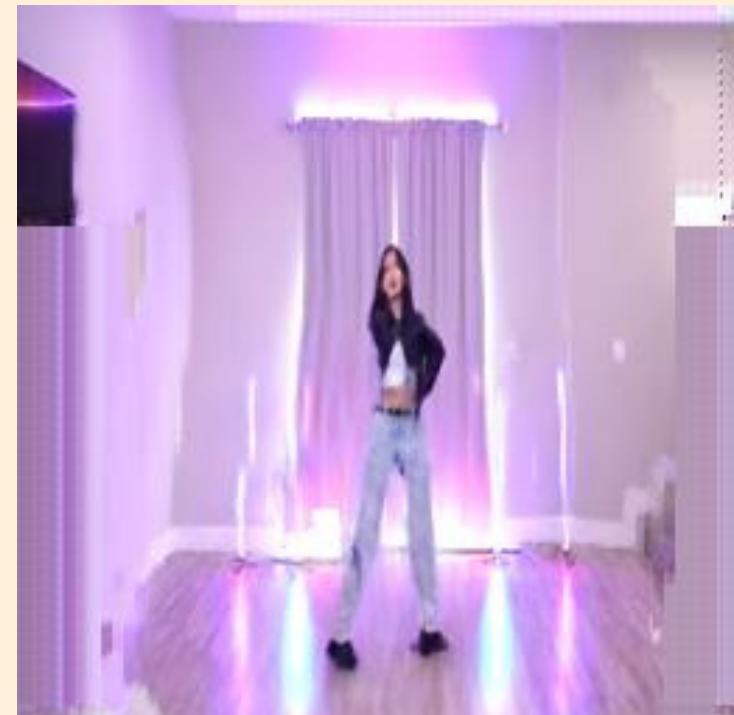
100



150



200 E



test data : 688.jpg  
optimizer : Adam

# References

## Our Github Repo

<https://github.com/Joannaaaaa/Synthesizing-Dynamic-Movements-for-Non-Experts>

## Contribution of each member

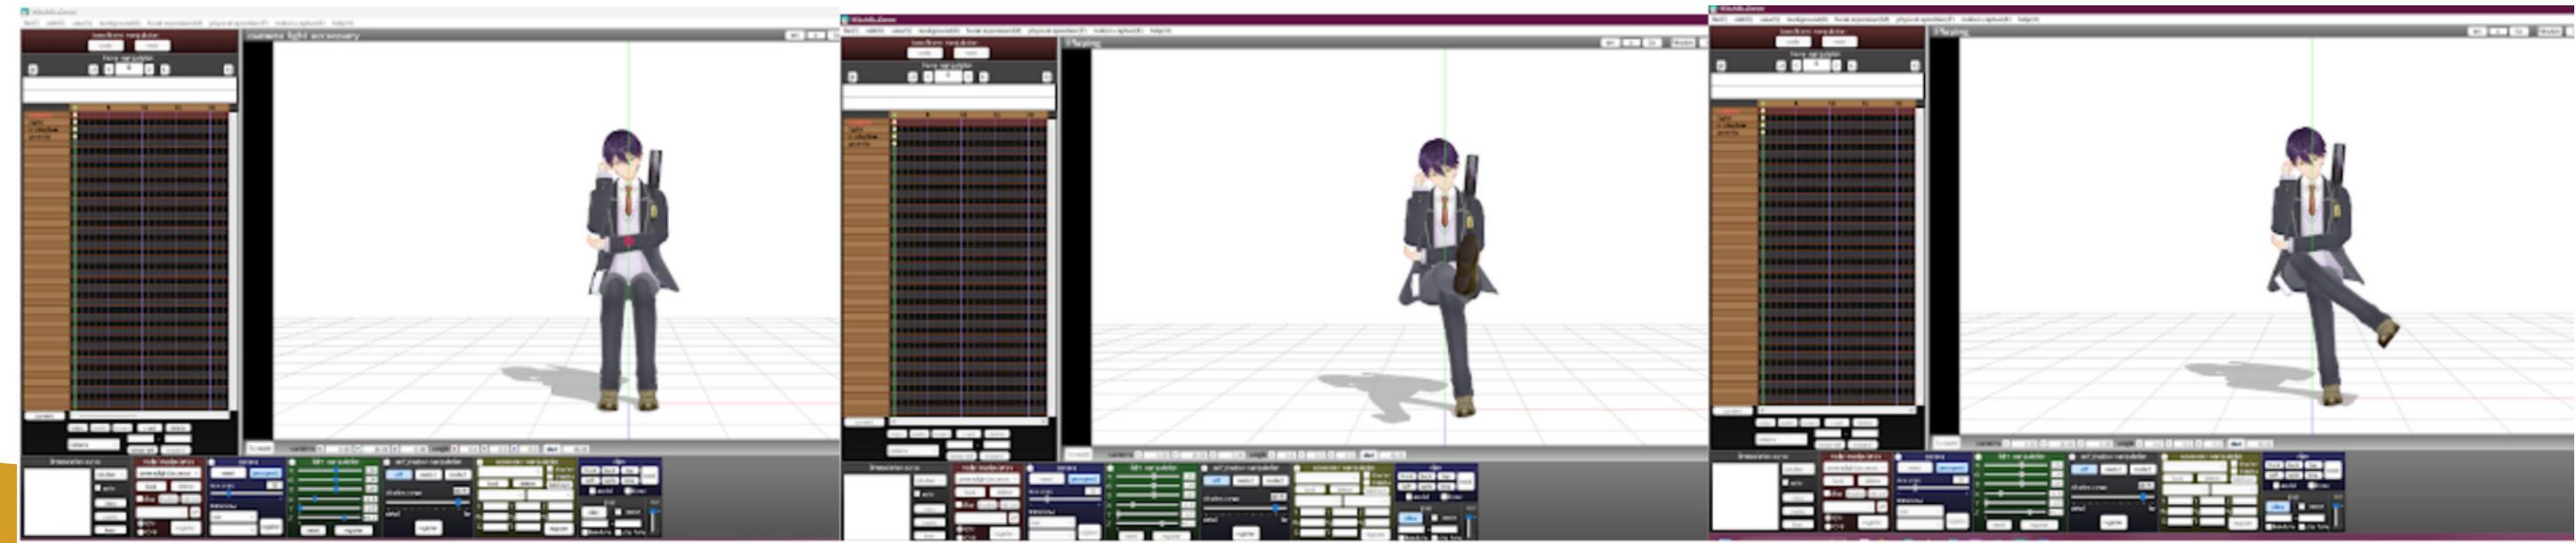
image to image - 楊芊華(35%)、賴怡暄(30%)

video to image、keypoint detection - 吳宜靜(35%)

# Future

## Tote bag with durable leather handle for heavier things

Presentations are tools that can be used as lectures, speeches, reports, and more. It is mostly presented before an audience. It serves a variety of purposes, making presentations powerful tools for convincing and teaching.



Thank You